



Fakultät für Informatik

Studiengang Informatik B.Sc.

Attack surfaces and security measures in enterprise-level Platform-as-a-Service solutions

Bachelor Thesis

von

Lukas Grams

Datum der Abgabe: TT.MM.JJJJ TODO

Erstprüfer: Prof. Dr. Reiner Hüttl

Zweitprüfer: Prof. Dr. Gerd Beneken

ERKLÄRUNG

Ich versichere, dass ich diese Arbeit selbständig angefertigt, nicht anderweitig für Prüfungszwecke vorgelegt, keine anderen als die angegebenen Quellen oder Hilfsmittel verwendet sowie wörtliche und sinngemäße Zitate als solche gekennzeichnet habe.

DECLARATION

I hereby declare that the thesis submitted is my own unaided work. All direct or indirect sources used are acknowledged as references. This paper was not previously presented to another examination board and has not been published.

Rosenheim, den XX.XX.2019 TODO

Lukas Grams

Abstract (german)

Die wachsende Zahl von Platform-as-a-Service (PaaS) Lösungen, Cloud-Umgebungen, containerisierter Datenverarbeitung und Microservice-Architekturen bieten neue Angriffsszenarien. Dies erhöht den Bedarf an neuen Verteidigungs-Strategien in der IT-Sicherheit von Entwicklungs- und Produktiv-Betriebsumgebungen. Gerade Lösungen auf Basis von (Kubernetes konformer) Container-Orchestrierung sind in der Wirtschaft stark gefragt und haben einen sehr unterschiedlichen Aufbau im Kontrast zu konventionellen und etablierten Lösungen. Dieser Umstand legt eine nähere Untersuchung dieses Teilbereichs nahe. Ziel dieser Arbeit ist es, folgende Fragestellungen zu beantworten:

- Was für Sicherheits-Risiken existieren für den Anbieter und/oder Nutzer einer mandantenfähigen PaaS-Lösung, wenn jeder Mandant über eigene Entwicklungs-, Verteilungs- und Laufzeit-Umgebungen für seine Anwendungen verfügt?
- Wie kann ein Anbieter von PaaS-Lösungen an interne und/oder externe Nutzer diese Risiken eindämmen?
- Was spricht in diesem Kontext aus Sicht eines PaaS-Anbieters jeweils für und gegen selbst betriebene bzw. Cloud-Lösungen?

Ein weiteres Ziel ist es, Maßnahmen für die unterschiedlichen Implementierungsmöglichkeiten zu empfehlen. Der Vergleich bestehender IT-Sicherheits-Risiken und daraus abgeleitet der priorisierte Handlungsbedarf sollen hierbei als zentraler Betrachtungswinkel dienen. Unter den etablierten PaaS-Lösungen in unterschiedlichen Umgebungen finden sich OpenShift Container Platform als Vor-Ort-Lösung und Azure Kubernetes Service für Cloud-Umgebungen. Um die genannten Ziele zu erreichen, grenzt diese Arbeit den Problembereich zuerst ein, indem sie sich auf die standardmäßig eingerichteten und zum Betrieb notwendigen Komponenten dieser zwei gängigen und Kubernetes-zertifizierten PaaS-Lösungen konzentriert.

Das Hauptaugenmerk liegt hierbei auf den Kubernetes-konformen Teilkomponenten, da hier die Risiken und Maßnahmen den weitreichenden Gültigkeitsbereich haben, unabhängig von der betrachteten Lösung. Aus den Zielen werden drei gängige Angriffsszenarien hergeleitet:

- Angriff von böswilligen Dritten auf die Infrastruktur ausgehend des LAN und/oder dem Internet
- Angriff von böswilligen Dritten aus einem Container heraus, über den die Kontrolle übernommen wurde. Ein Beispiel wäre das Ausführen von Code oder Befehlen per Zugriff von außen.
- Angriff von böswilligen Dritten sowie böswilliger oder fahrlässiger Nutzer bzw. Nutzer-Identität, was ein Kompromittierungsrisiko für diesen und/oder weitere Mandanten und deren Anwendungen darstellt.

Diese Szenarien dienen als Grundlage zur Identifikation von Angriffsvektoren, deren jeweiliges Schadenspotential evaluiert und das bestehende Risiko eingeschätzt werden. Es werden potentielle Maßnahmen zur Risiko-Eindämmung gesucht, evaluiert und exemplarisch in der Praxis umgesetzt. Im Anschluss wird eine Neubewertung des Risikos vorgenommen, um so beispielhaft eine zielführende Methode zum Risikomanagement zu demonstrieren. Mithilfe der Ergebnisse werden verschiedene Implementierungs-Empfehlungen verglichen, anhand der verschiedenen Anwendungsfälle differenziert und jeweils Maßnahmen empfohlen.

Schlagworte:

Platform-as-a-Service, PaaS, Cloud, Security, Container, Kubernetes, Docker, Risikomanagement, OpenShift Container Platform, OCP, Azure Kubernetes Service, AKS

Abstract (english)

The increasing amount of Platform-as-a-Service (PaaS) solutions, cloud-hosted environments, containerized workloads and microservice architectures introduce new attack scenarios. This creates the need for new defense strategies in both Development and Operations. Especially solutions providing (Kubernetes compliant) container orchestration are identifiably different and high in business demand compared to long established solutions. This calls for a more detailed, focused examination. The thesis aims to answer the following questions:

- What generic security risks emerge when providing or using a multi-tenant PaaS solution, with each tenant developing, deploying and running their own applications?
- How can a PaaS provider (serving internal and/or external users) mitigate those risks?
- In this scope and from a PaaS provider viewpoint, how does an on-premise solution compare to a public cloud solution?

Another goal is to recommend security measures for different implementation use cases. A comparison of existing IT security risks and the need for action derived from it should serve as central point of view. Examples for widely used PaaS solutions in different environments include OpenShift Container Platform as an on- premise solution and Azure Kubernetes Service as a public cloud solution. To achieve the aforementioned goals, the thesis will first limit the view on the problem to a manageable scope by concentrating on the components enabled by default and those required for operations of these established and Kubernetes compliant solutions. Components providing Kubernetes compliance will be the main focus, as these bear the most significance across all Kubernetes Certified solutions. Three common attack scenarios will be derived from the goals:

- Malicious third party attacking the underlying infrastructure from within the LAN and/or the internet
- Malicious third party attacking from inside a hijacked container, i.e. remotely executing code or commands

- Bad User, i.e. a negligent, hijacked or malicious developer account risking compromise of his own and/or other applications

These scenarios serve as a foundation to identify attack vectors, evaluate their respective potential impact and estimate their risks. Possible measures to mitigate those risks will also be explored, evaluated and exemplary put to use in practical examples. Subsequently the risk will be re-evaluated in order to illustrate a viable, result oriented risk management method. With the results gathered, the thesis will compare different best practice implementations for different use cases and recommend measures for each.

Keywords:

Platform-as-a-Service, PaaS, Cloud, Security, Container, Kubernetes, Docker, risk management, Open-Shift Container Platform, OCP, Azure Kubernetes Service, AKS

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Objective	1
1.3	Scope limitation	2
1.4	Research basis and its limits	3
2	Theory	5
2.1	Platform-as-a-Service	5
2.2	Containers	6
2.2.1	What are containers?	6
2.2.2	Differentiating docker	7
2.2.3	Images, image building and Dockerfiles	7
2.2.4	Container standards and interfaces	8
2.3	Container orchestration	8
2.3.1	Kubernetes	8
2.3.2	OpenShift	14
2.3.3	Azure Kubernetes Service	14
3	Deriving the attack surface and security measures	17
3.1	Defining procedure and approach	17
3.2	Identified vectors and security measures	18
3.2.1	V01 - Reconnaissance through interface components	18
3.2.2	V02 - Reading confidential information through interface components	19
3.2.3	V03 - Configuration manipulation through interface components	19
3.2.4	V04 - Compromise internal master components	20
3.2.5	V05 - Image poisoning and baiting	20
3.2.6	V06 - Configuration poisoning and baiting	21
3.2.7	V07 - Lateral movement through cluster	22
3.2.8	V08 - Container breakout	22
3.2.9	V09 - Image cache compromise	23

3.2.10	V10 - Container modification at runtime	23
3.2.11	V11 - Resource hoarding (sabotage)	23
3.2.12	V12 - Resource misuse (cryptojacking)	24
3.2.13	V13 - Adding rogue containers	24
3.2.14	V14 - Adding rogue nodes	25
3.2.15	V15 - Leveraging bad user practice	25
3.2.16	V16 - Leveraging bad infrastructure	25
3.2.17	V17 - Leveraging bad patch management	26
4	Assessing the attack surface risk	29
4.1	Defining procedures and approach	29
4.1.1	Specification of environment factors	29
4.1.2	Deriving a risk estimation formula	30
4.1.3	Resulting formula	33
4.2	Estimating the solution specific risk	33
4.2.1	V01 - Reconnaissance through interface components	37
4.2.2	V02 - Reading confidential information through interface components	37
4.2.3	V03 - Configuration manipulation through interface components	37
4.2.4	V04 - Compromise internal master components	38
4.2.5	V05 - Image poisoning and baiting	38
4.2.6	V06 - Configuration poisoning and baiting	39
4.2.7	V07 - Lateral movement through cluster	39
4.2.8	V08 - Container breakout	39
4.2.9	V09 - Image cache compromise	40
4.2.10	V10 - Container modification at runtime	40
4.2.11	V11 - Resource hoarding (sabotage)	40
4.2.12	V12 - Resource misuse (cryptojacking)	41
4.2.13	V13 - Adding rogue containers	41
4.2.14	V14 - Adding rogue nodes	42
4.2.15	V15 - Leveraging bad user practice	42
4.2.16	V16 - Leveraging bad infrastructure	42
4.2.17	V17 - Leveraging bad patch management	43
5	Managing the attack surface risk	45
5.1	Defining procedures and approach	45
5.2	Managing the risk of V07 - Lateral movement through cluster	45
5.2.1	Demonstrating the successful attack without security measures	45

5.2.2	Selecting security measures	46
5.2.3	Demonstration with implemented security measures	46
5.2.4	Risk reassessment	46
5.3	Managing the risk of V08 - Container breakout	46
5.3.1	Demonstrating the successful attack without security measures	46
5.3.2	Selecting security measures	46
5.3.3	Demonstration with implemented security measures	46
5.3.4	Risk reassessment revisited	46
5.4	Continuing the risk management process	47
6	Conclusion	49
6.0.1	Comparing on-premise and public cloud	49
6.0.2	Summary	49
A	Appendix	i
A.1	AKS versions available on May 3rd, 2019	i
A.2	Security advisory email from Microsoft	ii
A.3	AKS cluster default configs?	ii
	Bibliography	v

List of Figures

2.1	Comparison of responsibilities in different service models ¹	5
2.2	Comparison of different application deployments on the same hardware ²	7
2.3	An overview of different Kubernetes (k8s) components ³	10
2.4	Connection between deployments and other k8s objects down to containers ⁴	11
2.5	Connection between deployments and other k8s objects down to containers ⁵	12
A.1	List of k8s versions available in AKS during the practical part of the thesis on May 3rd, 2019. Screenshot taken by Lukas Grams.	i
A.2	An email from July 15th, 2019, advising users to reboot their AKS clusters in order to apply security patches. Screenshot taken by Lukas Grams.	ii
A.3	Default Azure Kubernetes Service (AKS) cluster configuration as of July 3rd, 2019. Screenshot taken by Lukas Grams.	iii

1 Wat17, section 'Original reference image'.

2 Aut19c, section 'Going back in time'.

3 Mei19.

4 Mei19.

5 Mei19.

List of Tables

4.1	A comparison of the resulting risk for the OCP and AKS environments	34
4.2	The risk estimation of all vectors for an OCP 3.11 cluster	35
4.3	The risk estimation of all vectors for an AKS 3.11 cluster	36

Listings

2.1 Exemplary .yaml file of a simple k8s deployment ¹	13
--	----

¹ Aut19a.

List of abbreviations

AKS	Azure Kubernetes Service
CIS	Center for Internet Security
CLI	Command Line Interface
CNCF	Cloud Native Computing Foundation
CRI	Container Runtime Interface
CSVS	Container Security Verification Standard
IaaS	Infrastructure-as-a-Service
k8s	Kubernetes
LXC	LinuX Containers
OCI	Open Container Initiative
OCP	OpenShift Container Platform
OKD	Origin Community Distribution of Kubernetes
OWASP	Open Web Application Security Project
PaaS	Platform-as-a-Service
RAL	required access level
RBAC	Role-based access control
RHEL	Red Hat Enterprise Linux
SCC	security context constraints

VM virtual machine

1 Introduction

With this chapter, the reader should be able to comprehend why this thesis was written, what it tries to accomplish and which topics are considered within the scope of this work.

1.1 Motivation

The increasing amount of Platform-as-a-Service (PaaS) solutions, cloud-hosted environments, containerized workloads and microservice architectures introduce new attack scenarios. This creates the need for new defense strategies in both Development and Operations. Especially solutions providing Kubernetes (k8s) compliant container orchestration are identifiably different and in high demand compared to long established solutions. This calls for a more detailed, focused examination.

1.2 Objective

This thesis aims to answer the following questions:

- What generic security risks emerge when providing or using a multi-tenant PaaS solution, with each tenant developing, deploying and running their own applications?
- How can a PaaS provider (serving internal and/or external users) mitigate those risks?
- In this scope and from a PaaS provider viewpoint, how does an on-premise solution compare to a public cloud solution?

Another goal is to recommend security measures for different implementation use cases. A comparison of existing risks and the need for action derived from it should serve as central point of view. Possible measures to mitigate those risks shall be explored, evaluated and exemplary put to use in practical examples. Subsequently the risk will be re-evaluated in order to illustrate a viable risk management method. With the results gathered, the thesis will compare different best practice implementations for different use cases and recommend measures for each.

1.3 Scope limitation

The thesis will first limit the view on the problem to a manageable scope by concentrating on the components enabled by default and those required for operations of two established and k8s compliant solutions. These solutions will be OpenShift Container Platform (OCP) as an on-premise solution and AKS as a public cloud solution. The latest stable version of OCP during the work on this thesis was version 3.11¹, which is based on v1.11 of k8s. Although k8s v1.13 was already available through AKS at the same point in time (see figure A.1), the available v1.11 was chosen to improve comparability.

The components of these two solutions providing k8s compliance will be the main focus, as these bear the most significance across all Kubernetes Certified solutions. A look at popular tools and frameworks used in such clusters will be avoided in order to keep the scope manageable, though some might be recommended as a mitigation. In order to be applicable to a higher number of use cases, attacks and measures seen in environments with exceptionally high security requirements might be mentioned, but not covered in their entirety. This includes state-sponsored actors deploying zero-day exploits, which are not applicable to a majority of solutions deployed and thus disregarded for the given context. This thesis aims to provide insight to the risks of providing a PaaS solution and mitigation thereof. As such, it will look at the capabilities a potential provider has to (mis-)configure such solutions - inherent risks of the technologies themselves are only explored when measures to mitigate them are accessible from a provider standpoint. In short, the goal is to improve the security of a given k8s cluster, not k8s itself. To follow the Open Web Application Security Project (OWASP) Risk Rating Methodology² down the line, threat actors have to be defined and grouped when applicable³. Examining a list of threat actors published through SANS⁴, only state-sponsored actors are excluded. This leaves us with cyber criminals, hacktivists, systems administrators, end users, executives and partners. Grouping the remaining actors by the factors used to estimate risks in section 4.2, three scenarios that encompass all of the actors can be identified:

- Malicious third party attacking the solution from within the company network and/or the internet
- Malicious third party attacking from inside a hijacked container, i.e. remotely executing code or commands
- Bad user, i.e. a negligent, hijacked or malicious account risking compromise of their own and/or other applications

1 Inc18a, see headline and date.

2 Fou19a.

3 Fou19b, Section 'Define all possible threats'.

4 Irw14, p. 12 to 17.

1.4 Research basis and its limits

During the research for this thesis, a considerable amount of sources has been examined. Surprisingly, very little has been found regarding a risk-based point of view on k8s solutions. A lot of sources start with measures and end with attacks, including one of the few published books¹. They recommend specific security measures and explain the sort of attacks they defend against. The only commonly referenced source that specifically introduces major risks for container technologies, without rating them, is the NIST Application Container Security Guide.² In addition to that, many of them are still work in progress or outdated.

Some of the sources commonly referenced include (in alphabetical order):

- The Center for Internet Security (CIS) Docker Benchmark³
- The CIS Kubernetes Benchmark⁴
- The book "Kubernetes Security - How to Build and Operate Applications Securely in Kubernetes"⁵
- The NIST Application Container Security Guide⁶
- The OWASP Container Security Verification Standard (CSVS)⁷
- The OWASP Docker Top 10⁸

Other sources include solution specific advisories, i.e. for k8s in general⁹ as well as OCP¹⁰ and AKS¹¹ specifically. In addition to that, recorded and openly available talks by different people in the industry are available. Especially the talks at KubeCon, the annual k8s conference, are accessible through the Cloud Native Computing Foundation (CNCF) YouTube presence¹². The speakers include recognized people in the industry, contributors to the k8s project and employees of reputable companies like Google, Red Hat and Microsoft. A considerable number of other online sources were referenced in the sources above or found with online search engines. These common recommendations exist, but neither claim to be exhaustive nor applicable to all future versions. In case of the CIS benchmarks, they are not

¹ RH18.

² SMS11.

³ Int19a.

⁴ Int19b.

⁵ RH18.

⁶ SMS11.

⁷ RF19.

⁸ Fou19c.

⁹ Aut19b.

¹⁰ Inc19a.

¹¹ Cor18.

¹² Channel Link: <https://www.youtube.com/channel/UCvqbFHwN-nwalWPjPUKpvTA/videos>

1 Introduction

even intended to be followed in full, but just as basis for security considerations.¹ In accordance to this, the findings can only be provided on a best-effort basis and not without reservation towards possible changes through future software versions or new information. This should be an information basis and approach for rating your own implementation while taking the business risk into consideration, not a definitive guide to all aspects in this regard. As with all systems, new attacks and vulnerabilities may emerge at any point in time. With sufficient research, the chance to have identified the most common attack vectors is probable but the limitations of this research has to be emphasized.

¹ McC18, starting at 29:18.

2 Theory

With this chapter, a reader with foundational knowledge of topics regarding Computer Science and/or Informatics will be able to grasp the specialized technologies discussed within the thesis and familiarize themselves with the definitions and terminology used throughout it.

2.1 Platform-as-a-Service

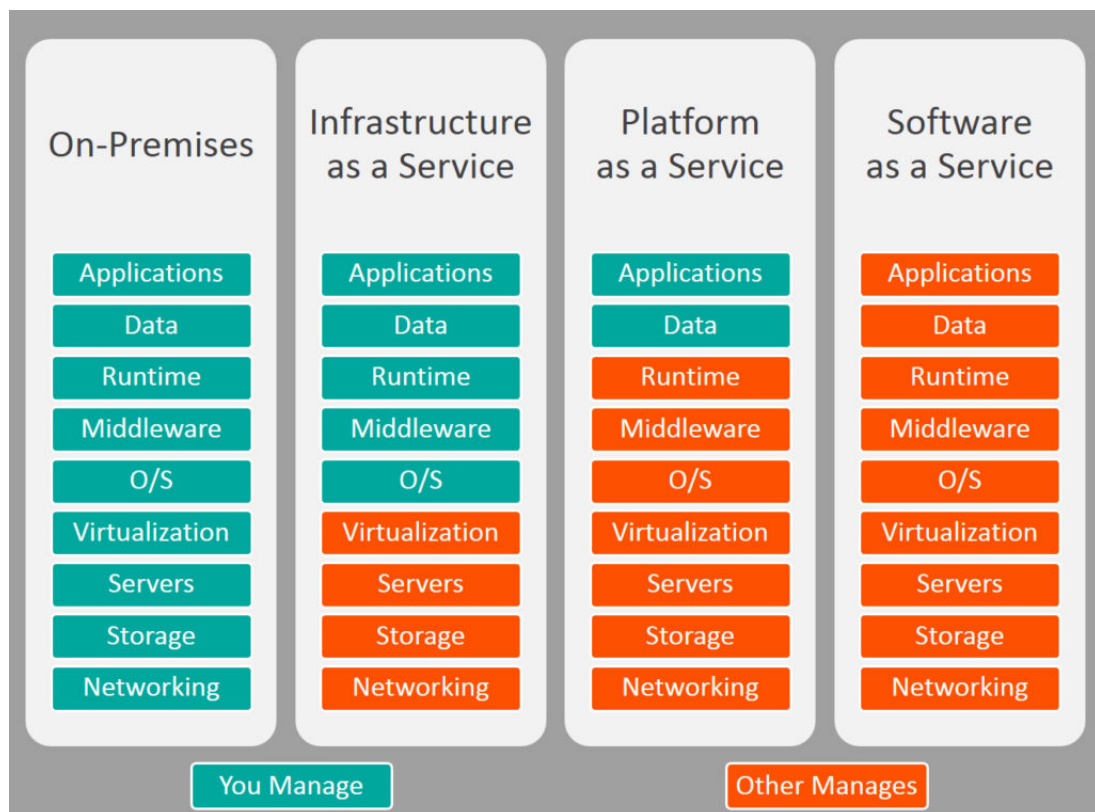


Figure 2.1 Comparison of responsibilities in different service models^a

^a Wat17, section 'Original reference image'.

2 Theory

In Platform-as-a-Service (PaaS) environments, a consumer delegates the management and control of resources needed to deploy his applications to his PaaS provider, beyond those covered by Infrastructure-as-a-Service (IaaS).¹ In an ideal scenario, this leads to a consumer not having to concern himself with the underlying network, hardware, servers, operating systems, storage or even common middleware like database management or log collection and analysis² and allows them to focus on other tasks, i.e. application development. As a downside to this, a consumer might only have limited control on, among other factors, the software installed on the provisioned machines. Although this shifts some of the responsibility burden towards the provider, the situation isn't as clear-cut as one might think. Figure 2.1 shows middleware and runtime as provided, but there is no clear standard on what capabilities or tools are included. A consumer might require capabilities which aren't provided or wishes to avoid provider lock-in through proprietary tools, resulting in some middleware responsibilities falling back to the consumer. A consumer might also have to extensively configure or modify the application-hosting environment for compliance or security purposes. Even some low-level tasks aren't completely managed, i.e. virtual machine (VM) reboots to apply security updates.³

2.2 Containers

Unsurprisingly, a basic building block of running containerized applications are containers. In order to run and manage containers, several components and systems are needed. The most important ones will be introduced here.

2.2.1 What are containers?

A container is a standard unit of software that packages up code and all its dependencies so the application runs quickly and reliably from one computing environment to another.⁴ From a technical viewpoint, a container is an isolated process running in the userspace of a host OS. The host system shares the kernel with all containers on it and might share other resources, but containers are run from container images which include any code and dependencies needed in order to make them independent of the infrastructure they are running on (except the kernel).⁵

Linux containers were widely popularized through Docker, a container system initially based on Linux Containers (LXC).⁶ Containers provide isolation of multiple applications running on the same machine

1 ST11, p. 2 to 3.

2 Cor19a, section 'Advantages of PaaS'.

3 Cor19b, section 'Process Linux node updates and reboots using kured'.

4 Inc19b, section 'Package Software into Standardized Units for Development, Shipment and Deployment'.

5 Raa18, slide 13.

6 Osn18, section '2013: Docker'.

and are often deployed in environments where this isolation would formerly have been achieved by using multiple VMs. Thus, they are often compared to them despite the fundamental technical differences. The Kubernetes documentation illustrates the differences as seen in Figure 2.2.

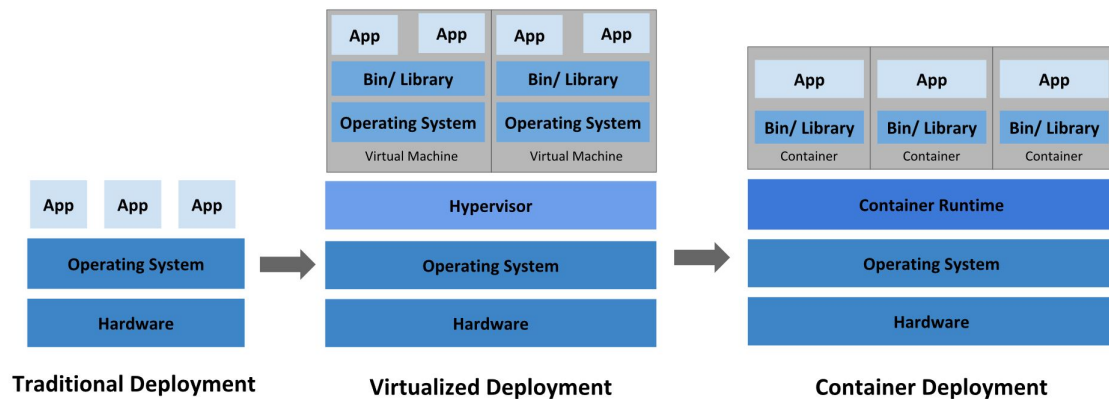


Figure 2.2 Comparison of different application deployments on the same hardware^a

^a Aut19c, section 'Going back in time'.

2.2.2 Differentiating docker

Talking about docker can be quite difficult, since the term is overloaded with different meanings - a company (Docker Inc.), their commercial products (Docker CE and Docker EE) and the former name of their open source project (formerly known as Docker, now called Moby).¹ Additionally, there is a Command Line Interface (CLI) called docker engine, which serves as an interface to the containers running on a host. It includes a high-level container runtime², which will be talked about in the section 2.2.4. Some sources also talk about docker-formatted containers when those technologies implement the same interfaces as the docker container runtime.³

2.2.3 Images, image building and Dockerfiles

A container image is a binary including all the data needed to run a container. It might also contain metadata on its needs and capabilities, i.e. version information through tags.⁴ Container images are sometimes referred to as docker images or docker-formatted images, but they can be run by other container runtimes and vice versa. In order to create a container image, you have to build it. This is often

¹ Cha17, section 'What is Moby?'

² Inc19c, section 'Develop, Ship and Run Any Application, Anywhere'.

³ Inc19d, section '1.11. Working with Docker formatted containers'.

⁴ Inc18b, section 'Docker Images'.

2 Theory

done through a build process executed by a container runtime. The instructions for container builds are commonly defined and documented in a Dockerfile¹ (which may also be done by non-docker programs, adding onto the vocabulary confusion). Container images can be distributed through container image registries, where images can be uploaded to and downloaded from. A commonly known example is docker hub, a public registry run by Docker Inc.

2.2.4 Container standards and interfaces

Without going into the nuances and historical developments, there are a multitude of programs mostly implementing three interfaces for container management. The two basic interfaces are the runtime and image specifications under the Open Container Initiative (OCI) which standardize how containers and container images should be formatted and executed.² The OCI also maintains a commonly used reference implementation called runc, alternatives include rkt and lxc.³ Runc and similar programs implementing these specifications are commonly called low level runtimes, in contrast to the high level runtimes that control them. These high level runtimes like containerd or CRI-O commonly manage more abstract features like downloading and verifying container images.⁴ Many high level runtimes today adhere to the Container Runtime Interface (CRI) so they can be used interchangeably by container orchestrators.⁵

2.3 Container orchestration

Once you want to use multiple containers on different machines talking to each other and offering stable services that continue even when a container or machine fails, the need for automated systems to manage these containers arises. Orchestrators can also provide other advantages like load balancing and automated scaling. Kubernetes systems are popular orchestrators currently in use. The sum of hosts running the orchestrator and containers are called a cluster.

2.3.1 Kubernetes

At its core, k8s is a control system for containers. It constantly compares the current state with the set target state and tries to correct towards the target when needed, i.e. when a container crashes. The

1 Inc18c, section 'Dockerfile reference'.

2 Fou19d, first paragraph.

3 Lew19a, section 'Examples of Low-Level Container Runtimes'.

4 Lew19b, Intro and section 'Examples of High-Level Runtimes'.

5 Aut19d, section 'Purpose'.

intended way for a user to deploy or change their application is to adjust the target state and let the k8s system take care of the rest.¹ There are many k8s distributions, many of which are certified Kubernetes offerings, meaning they all comply to the same standards and interfaces. These are set by the Linux Foundation through the CNCF which oversees the project. The Kubernetes code base is open source and maintained on GitHub, but any implementation fulfilling the (publicized) requirements can become k8s certified, regardless of how much code they changed.² You could look at k8s as a standardized interface for container orchestration with a public reference implementation.

Kubernetes components

In order to deliver a working k8s cluster, multiple (binary) components are needed. Master components provide the control plane, while node components are run on each underlying machine in order to maintain and provide the environment to execute the containers you want to run eventually. Master components are often exclusively run on machines dedicated to them, which are called master nodes - in contrast to worker nodes, which run the containers your applications consist of.³

The most relevant master components from the perspective of this thesis are:

- The kube-apiserver, which exposes the Kubernetes API. It is the front-end for the Kubernetes control plane; Cluster user or Administrator commands are typically directed at and processed by this component.
- Etcd, a distributed high-availability key value store where, among other things, secrets and authentication information are stored.

The important node components include:

- The kubelet, an agent running on each node in the cluster. It mostly monitors the state of any containers started by the k8s cluster. These are run through pods, a Kubernetes object designated to executing containers. The kubelet also interacts with the master components and reports on the monitoring data.
- The container runtime as depicted in figure 2.2, which is responsible for actually running containers. Examples include OCP using Docker while AKS uses Moby, but any implementation of the CRI is supported.⁴

1 Aut19a, section 'Understanding Kubernetes Objects'.

2 Fou19e, section 'There are over 80 Certified Kubernetes offerings.'

3 Aut19e, section 'Master Components'.

4 For additional information, refer to section 2.2.4

2 Theory

Figure 2.3 illustrates these components in context. It is to note that users in this illustration are the users of the applications running in the cluster; From a provider standpoint the users would be the people responsible for development and operations.

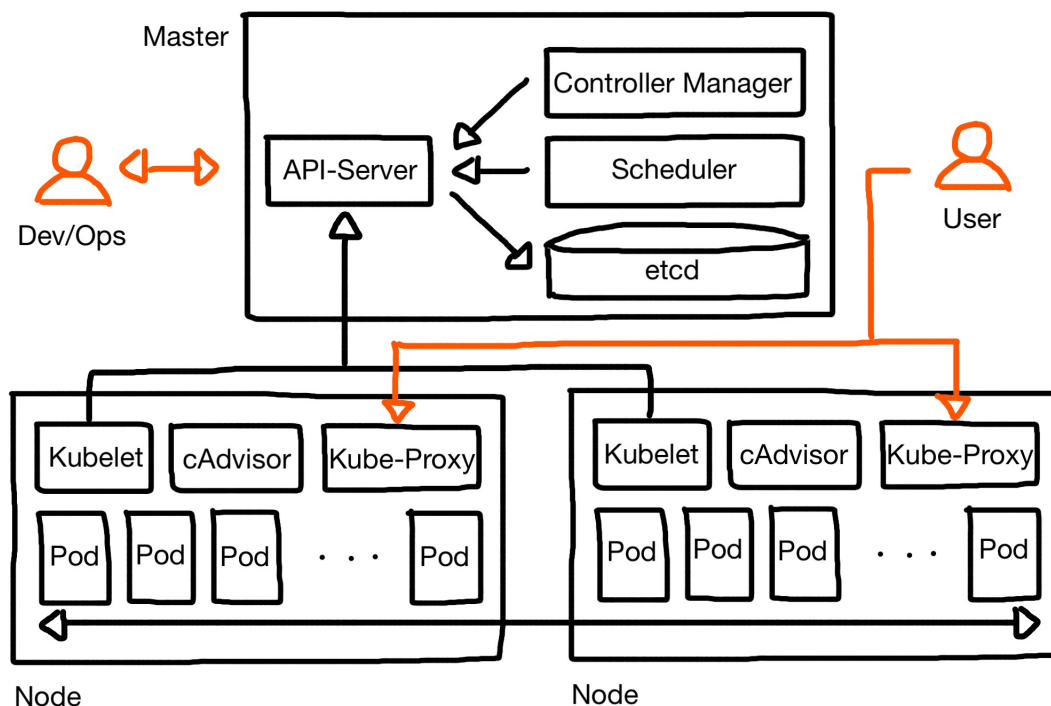


Figure 2.3 An overview of different k8s components^a

^a Mei19.

Kubernetes objects

Kubernetes objects are persistent entities that represent the desired state of your cluster. They can describe what containers should run, which resources are available to what system and the policies to apply (i.e. automatic restart behavior and communication restrictions). The intent is to modify these objects in order to change the target state, which the k8s system then works towards by adjusting the current state to match.¹ Some objects, i.e. pods, belong to a specific namespace, meaning a virtual cluster of many in a shared physical cluster. Others, like node objects describing the underlying machine, exist outside of a specific namespace. In order to understand how one can make the k8s system run an application according to its requirements, an understanding of the basic objects is needed.

A pod is the basic k8s object and encapsulates a container with some resources like an IP, storage and policies. Pods are typically each comprised of one container, a single instance of an application in k8s.

¹ Aut19a, section 'Understanding Kubernetes Objects'.




Figure 2.4 Connection between deployments and other k8s objects down to containers^a

^a Mei19.

They may contain more than one container for cases where these are tightly coupled and directly share resources.¹ That's all the pods do. They run.²

Pods are usually not created manually, but started and stopped by a replicaSet. Their purpose as a k8s controller is to maintain a stable set of replica pods running at any time in order to ensure availability of the function this type of pod provides.³

ReplicaSets are in turn managed by deployments. Just as replicaSets control pods, deployments control replicaSets in order to maintain the currently desired state of a cluster. Figure 2.4 illustrates this connection.

Since multiple identical pods may provide the same functionality and instances of this type of pod could be stopped or started at any point, how could a pod or other system address and connect to a pod? This can be solved by configuring a service, which abstracts a set of pods and their access policy. Typically, you talk to services instead of other pods directly. These services might then be published cluster-externally, i.e. through a load balancer provided by a cloud provider as seen in figure 2.5.

¹ Aut19f, section 'Understanding Pods'.

² BC18, p. 4.

³ Aut19g, introductory sentence.




Figure 2.5 Connection between deployments and other k8s objects down to containers^a

^a Mei19.

Using Kubernetes

There are many ways and solutions to set up a k8s cluster. Anyone can write their own one from scratch, but so-called turnkey solutions for cloud and on-premise environments exist, too, which significantly reduce the time and effort required to set up and run a cluster.¹ Once your cluster is set up, it is typically interacted with through the k8s API. For human interaction, kubectl is a CLI reference implementation which enables remote interaction with it.² In order to create or manipulate an object, you need to provide its new specification. This is typically supplied to kubectl as a .yaml file, an example of which can be seen in listing 2.1.³

Listing 2.1 Exemplary .yaml file of a simple k8s deployment⁴

```

1  apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5  spec:
6    selector:
7      matchLabels:
8        app: nginx
9    replicas: 2 # tells deployment to run 2 pods matching the template
10   template:
11     metadata:
12       labels:
13         app: nginx
14     spec:
15       containers:
16       - name: nginx
17         image: nginx:1.7.9
18         ports:
19       - containerPort: 80

```

¹ Aut19h, sections 'Turnkey Cloud Solutions' and 'On-Premises turnkey cloud solutions'.

² Aut19i, section 'The Kubernetes API'.

³ Aut19a, section 'Describing a Kubernetes Object'.

2.3.2 OpenShift

OpenShift is a commercial Kubernetes container platform by Red Hat. There are several different options to choose, deployable in different environments. OpenShift Online is hosted in a public cloud, while OpenShift Dedicated is hosted on dedicated nodes in a virtual private cloud. The option most relevant to this thesis is OCP, which can be deployed on any infrastructure, including on-premise environments.¹ Red Hat sponsors and supports the development of Fedora, a Linux distribution on which they base their Red Hat Enterprise Linux (RHEL) OS on. RHEL is then used by them commercially by supplying its binaries and updates through commercial licenses. Similar to this, Red Hat develops and supports a k8s distribution called Origin Community Distribution of Kubernetes (OKD), which then serves as a base for their OpenShift products, including OCP.² OKD release versions correspond to the k8s version it is based on, i.e. OKD 1.11 is based on k8s 1.11.³ OCP versions may differ, but Red Hat maintains a list of tested integrations for each major OCP release, giving insight to which k8s version it is based on.⁴ OCP is a certified Kubernetes distribution, which means it implements the same interfaces as all the others.⁵ Despite this, the differences in intended usage become apparent quite quickly, starting with the CLI used to interact with a cluster (`oc` instead of `kubectl`).⁶ Even basic concepts like namespaces are different, as OCP uses the similar but not identical concept of projects.⁷ Another difference of note is the mechanic of security context constraints (SCC), which are the OpenShift equivalent of and initial inspiration for k8s Pod Security Policies⁸. OCP uses Docker as its container runtime by default, but CRI-O is another option⁹.

2.3.3 Azure Kubernetes Service

The Azure Kubernetes Service is a managed k8s solution offered by Microsoft and running in the Azure public cloud. As it does not differ from the k8s reference implementation compared to OCP, its functionality is identical to the general k8s mechanics described in section 2.3.1. The computing resources used to run the cluster are pre-configured and provisioned VMs in the Azure cloud, which do not need to be configured and run a container runtime based on Moby, a toolkit from which the Docker runtime is built.¹⁰ k8s versions in AKS directly correspond to the k8s version they are based

¹ Inc19e, section 'OpenShift plans and pricing'.

² Har19, section 'OKD vs Red Hat OpenShift'.

³ Inc19f, section 'What is OKD?'.

⁴ Inc19g, table 'Platform Components'.

⁵ Fou19e, table 'Platform - Certified Kubernetes - Distribution'.

⁶ Inc19h, section 'Basic Setup and Login'.

⁷ Inc19i, section 'Overview'.

⁸ refer to section 3.2.5 for details on Pod Security Policies

⁹ Inc19j, section 'Getting CRI-O'.

¹⁰ Cha17, section 'What is Moby?'

on. As many cloud services do, there are multiple integrations to other cloud services, in case of AKS including an option to integrate your Azure Active Directory to authenticate users.¹

¹ Cor19c, first paragraph.

3 Deriving the attack surface and security measures

Picking up the three scenarios of section 1.3, the possible attack vectors are identified in this chapter to get an overview of the attack surface posed by k8s systems. Vectors and the possible attacks by different threat actors are explained as well as potential security measures introduced.

3.1 Defining procedure and approach

The identification of attack vectors within scope turned out to be a big challenge, both in research and preparation for a clear presentation. As new minor versions of k8s are released approximately every three months and the Kubernetes project only maintains release branches for the most recent three minor releases², the underlying system evolves continuously. Due to lengthy review processes, accredited literature about the topic are rare and risk being outdated quite quickly. Even less formal sources like blog posts or guidelines published by organizations are both rare, still unfinished at the point of this writing or updated continually, further complicating the creation of a snapshot regarding the current state. Nonetheless, it was attempted to achieve it on a best-effort basis. These vectors are applicable to all k8s solutions, although the available security measures will vary between different products or implementations. Thus, applicable generic measures will be mentioned and specific AKS and OCP examples may be provided for further illustration. In order to give a more comprehensible picture, the scope of these vectors varies. Some areas of interest, i.e. unpatched software and vulnerabilities in the underlying server infrastructure, have a broad scope. Since they are well known security topics applicable almost all information technology, they have been broadly encapsulated for a more clear understanding. If there would be rigorous differentiation between only slightly different vectors, this list would be multiple times as long as it already is.

² Aut19j, section 'Supported versions'.

3.2 Identified vectors and security measures

The following headlines will each respectively define and explain one of the 17 identified attack vectors of this thesis. For brevity of reference, each has a Vector ID assigned which is formatted as V_{xx}.

3.2.1 V01 - Reconnaissance through interface components

An attacker could gather information through the available interfaces in order to prepare further attacks and look for vulnerabilities to exploit. A k8s cluster has many interfaces, some of which are intended to be accessed by users or other non-master-components. These components include the kube-apiserver introduced in section 2.3.1, the popular web-based k8s user interface called Dashboard¹ as well as solution-specific interfaces like the OCP web console² or the Azure Portal for AKS³. Other interfaces are intended to be accessible from inside a cluster. These may include the components accessible from cluster-external sources, but also APIs for the integration of additional functionality. A common use case for such APIs is the provisioning of additional cloud resources in order to scale up the cluster. One example of reconnaissance against the kube-apiserver was demonstrated at KubeCon 2019⁴, while an example of leveraging the cloud api can be found in a blog post⁵.

Following the security principle of limiting the attack surface⁶, the measures against this include blocking access to or outright deactivating interfaces. If interfaces are unused, i.e. the k8s Dashboard where only CLI tools are used for cluster interaction, one can easily deactivate the Dashboard and thus reduce the attack surface without further work needed. If an interface is used, but alternatives exist, it may be advisable to evaluate whether a change in workflows or software architectures in order to decommission the interface may provide an overall benefit. Least privilege, another security principle⁷, can be followed by restricting access to used interfaces. Authentication can be required of both humans and automated processes and their access may be restricted to the information and capabilities strictly needed to operate. KubiScan, an open source tool to scan the Role-based access control (RBAC) permissions of a k8s cluster for permissions considered risky, is publicly available to assess a given cluster. Users might be required to log in through a two factor authentication to hinder access through stolen credentials. Network restrictions can be put in place to restrict the vantage points from which an attacker may gain access. Following Defense in Depth⁸, this should not only be done from the

1 Aut19k, first paragraph.

2 Inc19k, section 'Overview'.

3 Cor19d, section 'Create an AKS cluster'.

4 Ric19, starting at 8:45.

5 Gom18.

6 RH18, p. 4 to 5.

7 RH18, p. 3 to 4.

8 RH18, p. 3.

outside, but access from inside the cluster, too. As an example, most containers don't need access to k8s master components, so these requests can be blocked by k8s Egress Network Policies¹. The impact of successful attacks could be limited through Alerting and restriction of the attacks 'blast radius'. Logging and alerting processes may be implemented through k8s audit policies² for requests that are considered suspicious and do not occur in normal operations, examples of which can be found in online sources³. The 'blast radius' may be reduced by isolating different applications or teams through multiple namespaces and restricting access to non-namespaced resources to cluster administrators, thus isolating an attacker to a single namespace instead of the whole cluster. An even more drastic isolation could be achieved through isolation by cluster, resulting in separate k8s master components per isolated entity.

3.2.2 V02 - Reading confidential information through interface components

The components mentioned in section 3.2.1 could also be used by an attacker to gather confidential information, which can be either useful in and of itself or used to gain further access. Extracted private keys, tokens or passwords may be leveraged to authenticate with additional privileges in order to steal or manipulate container images, adjacent applications, source code and interact with other interfaces.

Despite presenting a more serious damage potential, the measures of section 3.2.1 also apply here since they target the same interfaces. In addition to that, special attention should be paid to the way secrets are passed into code running under k8s. Embedding secret information into container images should be avoided, since those are not only unencrypted and difficult to change, but may be extracted by an attacker reading the Dockerfile or inspecting an image themselves. Supplying them through environment variables is also not recommended, as those might be exposed in unconsidered ways. It is recommended to supply them to containers at runtime through k8s secrets or external tools like those integrated into cloud platforms or offered by third-party providers. Popular third-party solutions include HashiCorp Vault and CyberArk Conjur.⁴

3.2.3 V03 - Configuration manipulation through interface components

With write access to the components mentioned in section 3.2.1, an attacker may change configurations of the cluster or its environment, i.e. security controls set in place through the cloud provider platform. Examples include disabling security restrictions, rendering alerting processes dysfunctional by blocking their communication or swapping container images with functional, but compromised alternatives.

1 Aut19l, section 'The NetworkPolicy Resource'.

2 Aut19m, section 'Audit Policy'.

3 San17, section 'Alerting on the Kubernetes infrastructure'.

4 RH18, chapter 7.

3 Deriving the attack surface and security measures

As this has a higher potential for damage, special care should be taken. The measures of section 3.2.1 are applicable here, too. Special thought may be put into restricting write permissions to the most privileged user accounts, i.e. cluster or cloud administrators for actions that can influence systems outside of a namespace or cluster. In order to automatically deny actions, k8s Pod Security Policies¹ may be leveraged. Since these are non-namespaced k8s objects themselves, they could be circumvented by anyone with permissions to manipulate non-namespaced objects, i.e. cluster administrators.

3.2.4 V04 - Compromise internal master components

In contrast to the components mentioned in section 3.2.1, there are others which are only intended to be interacted with by other k8s master components. They may also be compromised by an attacker and are furthermore referred to as internal components. The most promising target is the etcd key value store introduced in section 2.3.1, as through it an attacker may gain any permission they want².

The available security measures are much more straightforward than those of the interface components, which is why all (ab)use cases have been united in one vector. Master components can be run on any machine in the cluster, but are often run on dedicated master nodes.³ Therefore, k8s solutions can restrict access to these by either not exposing them to anyone outside the pool of master nodes or, in case the components themselves are deployed as containers, restricting write access to their namespace they are running in as well as communication to the containers themselves. Synonymous to section 3.2.1, isolating through separate clusters per team or application can restrict the damage potential through separate master components.

3.2.5 V05 - Image poisoning and baiting

An attacker may try and lead a cluster user to run containers from a container image that is either controlled by the attacker or insecure.⁴ This could be done as an untargeted attack by offering images on popular public image repositories like Docker Hub or providing examples of similarly bad Dockerfiles in tutorials or support forums. Since building new containers from scratch is a lot of work, container images for popular applications are often sought to either build upon or use out-of-the-box. Applications for which no official container image is maintained are an especially common niche for such attacks. In contrast to this, an attacker could also try a more targeted attack by uploading bad images to the (possibly private) repository used by their target. Even previously trusted images may become insecure

¹ Aut19n, section 'What is a Pod Security Policy?'

² RH18, chapter 'Running etcd Safely'.

³ Aut19e, section 'Master Components'.

⁴ SMS11, p. 13 to 14.

over time as new vulnerabilities are publicized, thus becoming "stale".¹ It should be noted that providing Dockerfiles instead of images may be easier to identify as malicious, but when a user builds an image from it anyway, the build process will commonly run under a root user.²

In order to mitigate all of the above, only safe images should be used to run containers within a cluster. Although this sounds simple in theory, it may prove to be far more complex in practice. One would first have to define what constitutes a safe image in their context before implementing any measures ensuring their usage. Docker Hub introduced measures that help identifying more trustworthy images by labeling them as Verified Publisher Images or Certified Images respectively³. In order to restrict the amount of possible images used, a private registry could be used and the cluster restricted to only using images from that registry. User may then be allowed to build new images based on those in the registry and upload those to it. The damage potential of images may be limited by restricting the capabilities a container may gain through the Pod Security Policies introduced in section 3.2.3. Images could also go through a (either manual or automated) vetting process before being admitted to a registry. Images in a registry may be reassessed on a regular basis, too, so images with known vulnerabilities can be identified and removed. Multiple tools that provide automated image scans exist and can be found online⁴, some of which provide integration into k8s clusters or cloud platforms.

3.2.6 V06 - Configuration poisoning and baiting

An attacker may try and lead a cluster user or even administrator to configure the cluster in such a way that the attacker gains control, the cluster is left insecure, or both.⁵ This vector has a lot in common with the one described in section 3.2.5. A bad configuration might include bad containers, but does not have to, which is why we kept this as a separate vector. Attacks like this are far easier to do untargeted by offering tutorials and posts in support forums, but might also be conducted in a targeted way, similar to spear phishing attacks. The number of possibilities for specific bad configurations is near endless. Easy and effective misconfigurations may be achieved by swapping out popular image names with similar ones controlled by the attacker (i.e. 'nodejs' instead of 'node'), preconfiguring weak or attacker known default passwords, simply omitting or misconfiguring security measures so they do not work effectively (i.e. defining Network Policies, but not applying them) or making internal interfaces available to the public network.

1 SMS11, p. 14.

2 While k8s currently doesn't provide image building capabilities out-of-the-box, many k8s certified solutions implement this functionality.

3 MI18, sections 'Verified Publisher Images and Plugins' and 'Certified Images and Plugins'.

4 RH19, section 'Securing your container images'.

5 SMS11, p. 13 to 14.

3 Deriving the attack surface and security measures

As there are many potential ways for misconfiguration, there are also many measures to mitigate them. Raising user and administrator awareness for the risk of using a configuration supplied by outsiders is recommended. Continuously training them to provide the knowledge needed to thoroughly understand and assess these configurations may help in mitigating this problem, too. Free and helpful tools in assessing configurations are kubesecc.io, kube-bench¹ and kubeaudit². Defining guidelines and best practices for configurations could be done and enforced by regular manual reviews or automated checks through tools like the Open Policy Agent admission controls³ or k8Guard⁴. As bad configurations may include bad containers, the measures introduced in section 3.2.5 are applicable, too.

3.2.7 V07 - Lateral movement through cluster

Once an attacker gains access to a container, he may try to access more lucrative information, applications or cluster components. Ways to achieve this could be sniffing network traffic or scanning the network for other containers, hosts, services, apis or similar interfaces.

Applicable security measures include network isolation through Network Policies⁵ and mutual TLS authentication for network communication through tools like Istio⁶.

3.2.8 V08 - Container breakout

A popular phrase about container security is that "[c]ontainers do not contain"⁷. Once an attacker is able to execute commands in a container within the cluster, he may try influence components outside its isolation confinements. If an attacker successfully gains access to the underlying system, they might gain visibility into or control over any container running on it, including those of other applications. Additionally, they could gain access to internal cluster components like the kubelet⁸. Ways to achieve this include the invocation of syscalls, accessing mounted parts of the host file system and requesting additional capabilities. More ways to achieve this are being discussed at the time of this writing as the OWASP Docker Top 10 are worked on⁹ and openly demonstrated¹⁰.

1 <https://github.com/aquasecurity/kube-bench>

2 <https://github.com/Shopify/kubeaudit>

3 con19, section 'Wrap Up'.

4 <https://k8guard.github.io/>

5 Aut19o, section 'The NetworkPolicy Resource'.

6 Aut19p, section 'Mutual TLS authentication'.

7 Wal14, first headline.

8 refer to section 2.3.1 for additional details

9 Fou19f, dialogue by GitHub users 'drwetter', 'gramsimamsi' and 'wurstbrot'.

10CB19, section 'Breaking down the proof of concept'.

This can again be mitigated by restricting the capabilities and host access a given container is provided through the Pod Security Policies mentioned in section 3.2.3. "Sandboxed" containers may be leveraged here, too, which provide more isolation between a container and its host at the cost of performance. Popular implementations include gVisor¹ and Kata Containers².

3.2.9 V09 - Image cache compromise

If the container image cached by a container runtime on the underlying host system is swapped out by an attacker, another container controlled by the attacker might be started.

In order to mitigate this, container runtimes may support capabilities to check for the integrity of a container image. An example of such features is Docker Content Trust, which enables integrity verification at runtime³.

3.2.10 V10 - Container modification at runtime

Instead of starting bad containers as discussed in section 3.2.13, an attacker may try to modify it in order to suit their needs and support further attacks. This may be done by downloading and running additional binaries, establishing connections to Command and Control Servers or manipulating the software and files in place.

Measures against this build include limiting the capabilities a container has as discussed in section 3.2.3. Limiting network connections to cluster-external sources may mitigate ways to download additional software or exfiltrate data. This could be done via Egress Network Policies, which were introduced in section 3.2.1 Monitoring tools could be used to detect this, of which several commercial products are offered and often integrated into (cloud) provider platforms⁴.

3.2.11 V11 - Resource hoarding (sabotage)

An attacker with the ability to use resources or manipulate configuration may misuse or misconfigure them in order to disrupt the availability of specific applications or the cluster as a whole in DOS attacks. This may be done in a myriad of ways. Exhausting the available resources through fork bombs⁵ or intensive use of memory, RAM, network bandwidth, processing power or available network ports

1 <https://github.com/google/gvisor>

2 <https://katacontainers.io/>

3 Inc19l, section 'About Docker Content Trust (DCT)'.

4 TK18, slide 40 to 41.

5 RH18, chapter 'Fork Bombs and Resource-Based Attacks'.

3 Deriving the attack surface and security measures

may be a straightforward technique. They could also delete data and binaries, i.e. k8s objects and node programs and operating system components, wherever they have access to. If an attacker tries to achieve disruption for a long period of time, they could misconfigure the setup in a way that is either difficult to debug or reverse.

Measures against this include dedicated master nodes, which are a common practice anyway, in order to keep entities without administrator access from affecting the master nodes and master components running on them. Synonymous to section 3.2.1, isolating applications or teams through different namespaces or even clusters would reduce the 'blast radius' of a successful attack. Once this isolation is in place, resource quotas¹ could be set in place to constrain the total resource consumption of a namespace. Tools to monitor usage² could be set in place, as well as alerting mechanisms for unusually high resource consumption. Misconfiguration could be mitigated by the same logging and alerting measures introduced in section 3.2.1. Deployment management tools like Helm could also be used, which simplify persisting the k8s object states in a version-control system.

3.2.12 V12 - Resource misuse (cryptojacking)

Instead of sabotaging the cluster and/or applications, an attacker may use resources they can allocate to achieve monetary gain. This is often done by mining cryptocurrencies. The most commonly cited example for this is the incident of electric car manufacturer Tesla, where attackers used the control gained over their cloud infrastructure to run mining software in order to gain cryptocurrency with resources paid for by the company³.

The measures against this include those introduced in section 3.2.11. As this vector might be more difficult to detect if done well, a focus on detection may be of use. More advanced measures are presented by RedLock, the company which uncovered the Tesla incident, on their blog⁴.

3.2.13 V13 - Adding rogue containers

In contrast to sections 3.2.5 and 3.2.10, an attacker or malicious user capable of starting containers in the cluster may simply start their own malicious one. Since they could try to configure it in any way they want, this could help them in further attacks, leveraging vectors like those described in sections 3.2.7 or 3.2.8.

1 Aut19q, section 'Viewing and Setting Quotas'.

2 Aut19r, first paragraph.

3 Tea18, section 'The Latest Victim: Tesla'.

4 Tea18, section 'Preventing Such Compromises'.

The mitigations of section 3.2.10 mostly apply here, except for those against downloading additional binaries at runtime. An attacker could simply supply these within a Dockerfile or container image so they would already be installed. Against this, the detection and mitigation measures against bad images described in section 3.2.5 are applicable.

3.2.14 V14 - Adding rogue nodes

An attacker could try and add a host as a cluster node which is controlled by them. If successful, containers will be scheduled on the new node, allowing an attacker access to these containers in order to read, exfiltrate or manipulate data. They would also have control over the kubelet on this node, as described in section 3.2.8. In a sufficient time period, the chance of any container in the cluster running on the node at least once increases. This process could be sped up by manipulating the reporting data sent by the kubelet to the kube-scheduler, faking a lot of unused resources on it so more containers are scheduled.

In order to mitigate this, nodes would have to authenticate themselves to the apiserver and permissions to add new nodes should be limited to the highest possible user accounts like cluster administrators.

3.2.15 V15 - Leveraging bad user practice

An attacker may leverage bad user practice to gain access to or permissions within a cluster. This vector comprises user practices outside of the cluster that lead to risks within it. Examples include phishing for user accounts, gathering keys/tokens published to public code repositories, gathering passwords published in credential leaks or dumps, scouting specific software or container images used as well as gathering logs published with information valuable to an attacker. This could be done in a targeted way (i.e. specific Open Source Intelligence gathering or sending spearphishing emails) or untargeted by searching large repositories like Github for information formatted in a specific way. A practical example would include looking for published '.kube\config' files, where authentication information and kube-apiserver addresses are saved to¹. Gaining access through tokens in public logs has been demonstrated, too².

3.2.16 V16 - Leveraging bad infrastructure

Even if the cluster itself is secured properly, an attacker may gain access through the underlying infrastructure if it isn't. This includes the configuration of cloud infrastructure as well as on-premise

¹ Aut19s, section 'Define clusters, users, and contexts'.

² EdO19, section 'Results and notable findings'.

3 Deriving the attack surface and security measures

configurations. Servers may open insecure or unnecessary ports towards the public, as well as interfaces that allow someone any sort of control, i.e. allowing public access to the Docker remote API¹. Internally, an attacker may access data belonging to other applications by leveraging side channel attacks like Spectre and Meltdown².

Measures against this include well known and documented security measures and best practices for traditional server infrastructure, including the full range of CIS benchmarks available³. Keeping the OS and programs of the underlying nodes minimal by only installing the necessary tools, functions and binaries may aid by reducing potential vulnerabilities on top of reducing the resource overhead. Establishing a 'bastion host'⁴ as single point of remote entry to the cluster may be advisable if such functionality is required. Measures can also be found in the OWASP CSVS⁵, CIS Docker Benchmark⁶, the CIS Kubernetes Benchmark⁷ and the NIST publication⁸. Cloud configuration may be assessed and/or enforced with automated tools like ScoutSuite⁹, cloud-custodian¹⁰ or environment-specific tools like azurite¹¹.

3.2.17 V17 - Leveraging bad patch management

An attacker may leverage publicly known vulnerabilities in software that is not up to date on any component used in the environment. This may even include systems assumed to be responsible for by other parties. Even in PaaS environments like AKS, servers have to be restarted by the cloud customer in order for some security updates to take effect. An example for this can be seen in an email sent to an Azure cloud account holder, a screenshot of which is provided in figure A.2.

Mitigating this is straightforward in theory - apply security updates, fixes and intermediary workarounds whenever they become available. In practice this proves to be an ongoing real-world problem, as the WannaCry epidemic very publicly demonstrated¹². Any and all unpatched components could pose a risk. Thus, responsibilities and guidelines should be defined for when and what to patch. Fallback plans accounting for vacation time and sick days should be implemented, too. The employees responsible for updating should set up or subscribe to their relevant vulnerability notification feeds, i.e. alerts

1 SN19, section 'Publicly Accessible Docker Hosts'.

2 Tec19, section 'Which cloud providers are affected by Meltdown?'

3 Int19c, refer to the list presented.

4 Sco19, section 'What is a bastion host, and do I need one?'

5 RF19, section 'Infrastructure'.

6 Int19a, chapters 1 through 3 and 5.

7 Int19b, chapters 'Worker Node Security Configuration' and 'Configuration Files'.

8 SMS11, chapters 3.5, 4.5 and 4.6.

9 <https://github.com/nccgroup/ScoutSuite>

10 <https://github.com/cloud-custodian/cloud-custodian>

11 <https://github.com/mwrlabs/Azurite>

12 Ked19, p. 2.

for new vulnerabilities¹. Servers should follow conventional patch management processes, as these are similar to conventional system infrastructure. Kubernetes distributions should be kept up to date with updates provided by the k8s distributor. Of note is the short time of guaranteed support by the k8s reference implementation. The Kubernetes project maintains the most recent three minor releases with a new minor version planned approximately every three months². This would result in a cluster upgrade to be required every nine months in order to stay supported. Other distributions may have longer time periods of guaranteed support, i.e. OCP³. Applying updates to containers by updating container images should be done regularly. This includes both incrementing the version numbers of their base images to a version without currently known vulnerabilities as well as updating binaries directly downloaded during the build process. As an 'update' command run by the package manager during the build process works dynamically, security updates to those packages are included any time an image is built again. In consequence, images should be rebuilt periodically. Using the image version tag 'latest' is not recommended, as this could automatically propagate bad images pushed to the image repository through to production environments⁴. Automated scanning may be implemented by open source tools like clair⁵. Limiting the probability for new vulnerabilities by reducing the software installed as described in section 3.2.16 may help, too.

1 Dio19, first paragraph.

2 Aut19j, section 'Supported versions'.

3 Inc19m, section 'OpenShift Container Platform v3'.

4 RH18, p. 40.

5 <https://github.com/coreos/clair>

4 Assessing the attack surface risk

With the attack vectors identified, we introduce a customized risk estimation model for our purpose and explain the challenges of developing it. The model is then used to estimate the relative risk of each vector, specific to each solution.

4.1 Defining procedures and approach

Finding a framework to properly conduct a risk assessment has proven to be more difficult and time consuming than expected. The process and result are laid out in this section.

4.1.1 Specification of environment factors

In order to achieve a view on the risks more accurately resembling situations where a solution might be implemented, several assumptions are made:

- People in contact with the solution are familiar with conventional security principles and measures, but are new to the technologies used in the solutions within scope. They might even be new to container and cloud solutions in general. This includes users of the solution like developers and project managers as well as operators and administrators.
- It is assumed that no special requirements like industry-specific compliance requirements have to be followed and the workloads processed within the solution have no exceptionally high security requirements.
- Some risks increase or decrease drastically, depending on many specific configurations. Considering the high system complexity, “getting it to work” is hard enough for users and operators new to these technologies.² When setting up and configuring a setup, the default configurations are left as-is whenever possible. Guidelines of specific implementations are followed, but whenever measures are presented as optional and not required, they will be skipped. Before recommending

² Gee17, starting at 3:05.

4 Assessing the attack surface risk

security measures, the setup will be modified just enough to become functional, without regards to the security implementations.

- Multiple tenants like different customers, teams or projects are separated by k8s namespaces or OCP projects respectively, not by isolated clusters.

In addition to these assumptions, two specific implementations were selected and respective software versions were chosen. Additional solutions or software versions have not been assessed due to the exponential increase in workload to do so, but the process outlined in this thesis could be used to assess them, too. The solutions were selected since they are both popular as well as there being specific interest in them by the company accompanying this thesis, HvS Consulting. These include AKS as a public cloud solution and OCP as an on-premise solution. As the latest stable version of OCP during the beginning of this thesis was v3.11, this version was selected. It is based on OKD v3.11, which in turn is based on k8s v1.11¹. AKS v1.13, directly correlating to k8s v1.13 has already been available at this point², but v1.11 was chosen in order to achieve better comparability between the two solutions. This thesis will not look further into the upsides and downsides of faster availability in contrast to longer release and support cycles.

Whenever possible, all default settings were used to set up both solutions and get them functional. As a managed solution, the AKS settings are less distributed and detailed. An overview of the default options at the time of preparing practical parts of this thesis can be seen in figure A.3. Setting up the OCP cluster was less straightforward and included steps with multiple options, but no defaults. The official setup guideline³ has been followed. A single master, single node setup was chosen. Multiple masters and (worker) nodes are recommended for production environments, but since this setup is purely used for assessment and demonstration, it is sufficient. Both nodes are virtual machines with sufficient resources in accordance to the documentation, running on RHEL 7.5. Best practices would recommend minimal OS setups like the RHEL Atomic Host offered by Red Hat, but this option has been chosen as such hosts are a common request within larger enterprises and as such faster and more easily available through preexisting processes in many corporate environments.

4.1.2 Deriving a risk estimation formula

Generic methods of estimating risk may not lead to the most accurate results when assessing specialized systems like orchestrated container environments. But since an initial search for risk assessment frameworks and methodologies specific to such environments proved unsuccessful, more generic

¹ Inc19g, refer to the second table below the headline 'Platform Components'.

² for more details, check figure A.1

³ Inc19n, sections 'Planning your installation' to 'Installing OpenShift'.

frameworks were considered. The process laid out in the OWASP Risk Rating Methodology¹ was an obvious choice, but discarded for the considerable effort needed to both map the generic ratings to indicators specific to our environment as well as estimating the sum of twelve to thirteen applicable factors per vector. Business impact factors would have been excluded since k8s solutions are assessed in general, but for the exception of financial impact in case of cryptojacking where financial damage can be detached from the applications running on it.

Another methodology considered but decided against was the Common Vulnerability Scoring System in its most current version, v3. As specifically stated in its user guide, it is not intended to measure risk but instead measuring the severity of specific vulnerabilities².

As an alternative, another methodology was proposed by the company accompanying this thesis, HvS Consulting. The generally accepted base formula of risk being the mathematical product of likelihood and impact forms the basis of this methodology, where the impact is estimated directly while the likelihood is calculated by summing up the estimation of several contributing factors. This is a simplified, but similar version of the OWASP Methodology³.

The original factors used to estimate a value for the likelihood were prevalence, awareness and exploitability. This methodology was customized in order to be more applicable, since prevalence would have been difficult to estimate and compare for the relatively new technologies. Instead of prevalence, the factors 'vantage point' and 'required access level' were added in order to explicitly consider both the possible origin of attack vectors as well as the privileges needed to leverage them. In contrast to other estimations they shall not influence the exploitability value, as the too much weight may then be put on these factors. The average value of these four factors lead to the resulting probability value. Customizing the factors is done to tailor towards our specific use case, a practice encouraged by the OWASP Risk Rating Methodology⁴.

As a simple comparison of the risk to each vector in relation to each other is sufficient to our use case, numerical values without unit instead of percentages or monetary amounts suffice. Compared to each other, they can be used to provide a prioritized list after the assessment. The values for each respective factor were developed with the three scenarios derived in section 1.3 in mind. The initial idea was to map the numbers of one through three to reasonable indicators for each factor with a value of three indicating the most serious contribution to overall risk. During the process, this was expanded to the values zero through four for severe outliers.

1 Fou19a, sections 'Step 2' to 'Step 4'.

2 FIR19, section '2.1. CVSS Measures Severity, not Risk'.

3 Fou19a, section 'Step 4: Determining the Severity of the Risk'.

4 Fou19a, section 'Step 6: Customizing the Risk Rating Model'.

4 *Assessing the attack surface risk*

The factor 'vantage point' increases in severity and value when it can be achieved more easily, as it would be easier to achieve it. It is valued at zero for situations that require more than a security baseline, i.e. physical access to the cluster nodes or hypervisors they may run on. A value of one is assigned for situations where commands need to be executed on the cluster nodes directly, as opposed to within the container. Command access to the containers is valued at two. Depending on the situation, access to the company network may be more difficult to achieve for an outside attacker than access to any single container, but from a cluster security perspective more people already have this sort of access. Therefore it is valued as a three, in contrast to the public internet accessible by anyone, which is an outlier in severity and assigned a value of four.

The 'required access level (RAL)' is incrementally increased when less privileges are needed, as this increases the factors contribution to the total risk. Administrator access to the cloud environment or infrastructure underlying cluster nodes (i.e. equivalent access to the hypervisor in case the cluster nodes are VMs) is valued at zero, as this would again exceed baseline security requirements. Administrator access to the orchestrator or nodes is valued at one. User access to these systems is valued at a two where read and write privileges are needed and valued at three in cases where only read privileges are required. If no privileges or authentication is needed at all, a sever outlier value of four is assigned.

'Awareness' represents how easy it is for an attacker to become aware of this vector being feasible in general and in a specific system targeted for attack. As this is not applicable to negligent users, the value represents how easy it is to accidentally 'leverage' this vector instead. The assigned value increases with increase possibility for awareness, increasing the total risk estimation. In order to be estimated to be unlikely and valued at one, custom-made tools would be needed to detect vulnerabilities or and accidental attacks could happen during uncommon tasks, both being custom to a system-specific environment. It is estimated to be possible and valued at two in cases where specific attacks could be done regardless of the specific environment but customization of generic tools is needed. If and accidental attack might happen during common tasks where tools or processes are not used the way they are intended, a value of two is also assigned. The highest awareness value of three is assigned when attacks are both environment-agnostic as well as detectable by readily available tools or by simply using graphical user interfaces. It is valued at three in cases where proper usage of tools or processes during common tasks might accidentally lead to a security problem.

The 'exploitability' factor differentiates between different levels of skill needed, increasing in value when less skill is required. In case of negligent users, it is simply valued the same as the awareness factor. It is valued at zero for situations that require more than a security baseline like unpublished and unknown exploitation of vulnerabilities. Synonymous to the definitions in the awareness factor, a difficult exploitability with a value of one would need custom tools for environment-specific exploitation, while an environment-agnostic exploit with require tool customization is valued as two. A value of three

is again assigned when attacks are environment-agnostic and exploitable without tool customization or simple user interface interaction.

The impact value was roughly estimated by differentiating between intermediate steps and achieved damage as well as whether a single namespace or the whole cluster is affected. Intermediate steps were assigned a value of one, while outcomes confined to a single namespace and limited monetary damages are given a value of two. Compromises of the whole cluster and potentially unlimited monetary damages are valued at three for the most severe outcome rating.

4.1.3 Resulting formula

The work elaborated on in the previous section leads to the following formula for estimating the total risk of the identified vectors, usable for a specific system setup at hand:

$$Total\ risk = Max[\frac{vantage\ point + required\ access\ level + awareness + exploitability}{4} * impact, 10]$$

This leads to a total risk of 0 through 10.5, which is rounded to full integers and capped at 10. In accordance to the OWASP Risk Rating Methodology¹, total risk values ≤ 3 are defined as low, ≤ 6 as medium and values above that are defined as high with the exception of 10, which is critical. It should be noted that the worst case option of each vector should be taken in case of multiple possible scenarios, attacks and outcomes.

4.2 Estimating the solution specific risk

The resulting risk estimation for both the OCP and AKS example systems are illustrated in Table 4.1. The individual values these results are derived from can be seen in tables 4.2 and 4.3. The details on how these values were determined are explained hereafter.

¹ Fou19a, section 'Step 4: Determining the Severity of the Risk'.

4 Assessing the attack surface risk

<u>Vector ID</u>	<u>Vector</u>	<u>OCP Risk</u>	<u>AKS Risk</u>
V01	Reconnaissance through interface components	3	3
V02	Reading confidential information through interface components	6	7
V03	Configuration manipulation through interface components	7	8
V04	Compromise internal master components	5	6
V05	Image poisoning and baiting	7	7
V06	Configuration poisoning and baiting	10	10
V07	Lateral movement through cluster	7	7
V08	Container breakout	5	7
V09	Image cache compromise	3	3
V10	Container modification at runtime	5	5
V11	Resource hoarding (sabotage)	8	6
V12	Resource misuse (cryptojacking)	5	8
V13	Adding rogue containers	5	6
V14	Adding rogue nodes	6	6
V15	Leveraging bad user practice	6	6
V16	Leveraging bad infrastructure	9	9
V17	Leveraging bad patch management	10	10

Table 4.1 A comparison of the resulting risk for the OCP and AKS environments

Vector ID	Vector	Vantage Point	RAL	Awareness	Exploitability	Probability	Impact	Resulting risk
V01	Reconnaissance through interface components	3	3	3	2	2.75	1	3
V02	Reading confidential information through interface components	3	3	3	3	3	2	6
V03	Configuration manipulation through interface components	3	1	3	2	2.25	3	7
V04	Compromise internal master components	3	1	3	0	1.75	3	5
V05	Image poisoning and baiting	4	4	3	2	3.25	2	7
V06	Configuration poisoning and baiting	4	4	3	2	3.25	3	10
V07	Lateral movement through cluster	2	2	3	2	2.25	3	7
V08	Container breakout	2	2	3	0	1.75	3	5
V09	Image cache compromise	1	1	2	2	1.5	2	3
V10	Container modification at runtime	2	2	3	2	2.25	2	5
V11	Resource hoarding (sabotage)	3	2	3	2	2.5	3	8
V12	Resource misuse (cryptojacking)	3	2	3	2	2.5	2	5
V13	Adding rogue containers	3	2	3	2	2.5	2	5
V14	Adding rogue nodes	3	1	2	2	2	3	6
V15	Leveraging bad user practice	3	4	2	2	2.75	2	6
V16	Leveraging bad infrastructure	4	4	2	2	3	3	9
V17	Leveraging bad patch management	4	4	3	2	3.25	3	10

Table 4.2 The risk estimation of all vectors for an OCP 3.11 cluster

Vector ID	Vector	Vantage Point	RAL	Awareness	Exploitability	Probability	Impact	Resulting risk
V01	Reconnaissance through interface components	4	3	3	2	3	1	3
V02	Reading confidential information through interface components	4	3	3	3	3.25	2	7
V03	Configuration manipulation through interface components	4	1	3	2	2.5	3	8
V04	Compromise internal master components	4	1	3	0	2	3	6
V05	Image poisoning and baiting	4	4	3	2	3.25	2	7
V06	Configuration poisoning and baiting	4	4	3	2	3.25	3	10
V07	Lateral movement through cluster	2	2	3	2	2.25	3	7
V08	Container breakout	2	2	3	2	2.25	3	7
V09	Image cache compromise	1	1	2	2	1.5	2	3
V10	Container modification at runtime	2	2	3	2	2.25	2	5
V11	Resource hoarding (sabotage)	4	2	3	2	2.75	2	6
V12	Resource misuse (cryptojacking)	4	2	3	2	2.75	3	8
V13	Adding rogue containers	4	2	3	2	2.75	2	6
V14	Adding rogue nodes	4	1	1	2	2	3	6
V15	Leveraging bad user practice	4	4	2	2	3	2	6
V16	Leveraging bad infrastructure	4	4	2	2	3	3	9
V17	Leveraging bad patch management	4	4	3	2	3.25	3	10

Table 4.3 The risk estimation of all vectors for an AKS 3.11 cluster

4.2.1 V01 - Reconnaissance through interface components

In the OCP setup, a user with read access to the apiserver or webinterfaces can scout out information. It is assumed that these are only accessible internally in typical corporate environments (Vantage Point: 3). By default, each account (project admin or project user, but not cluster admin) can only see information about their own project while a cluster admin can see all namespaces. The worst case assumption would be leveraging user permissions (RAL: 3). The information gathering processes and interfaces are known and documented pretty well (Awareness: 3), but the information gathered has to be analyzed specific to the environment (Exploitability: 2). This could show information useful to an attacker like software versions, running systems or pods, user account privileges or misconfigurations and may help in planning and confirming the effectiveness of further attacks (Impact: 1). With a total risk value of 3, it is classified as a low risk.

The circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet (Vantage Point: 4, RAL: 3, Awareness: 3, Exploitability: 2, Impact: 1). The total risk value remains unchanged.

4.2.2 V02 - Reading confidential information through interface components

In the OCP setup, a user with access to the apiserver or webinterfaces (Vantage Point: 3) and read access (RAL: 3) can gather confidential secrets like certs, tokens or passwords which are intended to be used by automated systems and/or users to authenticate themselves to cluster components and gain authorized access. They could pull/push images, trigger actions in other applications or containers and more. As a leakage of confidential information, this would constitute a security incident by itself (Impact: 2). Kube-hunter is a readily available tool that checks for some confidential information automatically (Exploitability: 3). With a total risk value of 6, it is classified as a medium risk.

Synonymous to V01, the circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet (Vantage Point: 4, RAL: 3, Awareness: 3, Exploitability: 2, Impact: 1). The total risk value is increased from 6 to 7 by this, classifying it as a high risk.

4.2.3 V03 - Configuration manipulation through interface components

In the OCP setup, administrators (RAL: 1) with access to the apiserver or webinterface (Vantage Point: 3) can change configurations on the cluster. This includes project-specific resources like pods, services, routes, but also cluster-global resources like nodes or authorization configuration (Impact: 3). The capabilities can be looked up through the documented console commands of kubectl (Awareness: 3), but ways to achieve long-term benefit as an attacker has to be analyzed environment-specifically

4 Assessing the attack surface risk

(Exploitability: 2). With a total risk value of 7, it is classified as a high risk.

Synonymous to V01, the circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet (Vantage Point: 4, RAL: 1, Awareness: 3, Exploitability: 2, Impact: 3). The total risk value is increased from 7 to 8 by this, classifying it as a high risk.

4.2.4 V04 - Compromise internal master components

In the OCP setup, misconfiguration of internal k8s components could lead to a full cluster compromise (Impact: 3). It is assumed that these are only accessible internally in typical corporate environments (Vantage Point: 3). The cluster configuration and all secrets and authorization credentials are stored in the etcd instance(s). This is secured by the OCP setup process and an administrator would have to willingly misconfigure the cluster (RAL: 1) through internal settings not intended to be changed in order to make this possible. Undisclosed vulnerabilities in the OCP could lead to this problem, too (Exploitability: 0). Kube-hunter checks for generally common cluster misconfigurations in this regard (Awareness: 3). With a total risk value of 5, it is classified as a medium risk.

Synonymous to V01, the circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet and internal component configuration being maintained by Microsoft. (Vantage Point: 4, RAL: 1, Awareness: 3, Exploitability: 0, Impact: 3). The total risk value is increased from 5 to 6 by this, still classifying it as a medium risk.

4.2.5 V05 - Image poisoning and baiting

In the OCP setup, the untargeted version¹ needs the least access. A free and anonymously created dockerhub account to upload malicious images is available from anywhere (Vantage Point: 4) by anyone (RAL: 4). By itself only a single application or namespace / project would be affected by this (Impact: 2), as further exploitation of this foothold for an attacker is handled in vector V07. These methods are publicly known and both the docker container runtime and docker hub actively try to mitigate this, but malicious images are only deleted when reported by enough users (Awareness: 3). Base containers and malware as well as known vulnerable versions are readily available from public sources, but need some technical expertise to 'weaponize' (Exploitability: 2). With a total risk value of 7, it is classified as a high risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

¹ refer to section 3.2.5 for a detailed description

4.2.6 V06 - Configuration poisoning and baiting

In the OCP setup, the untargeted version also needs the least access. Synonymous to V05, this can be done by anyone (RAL: 4) from anywhere (Vantage Point: 4). Compromised configurations can lead to a full cluster compromise (Impact: 3). Examples are readily available from public sources (Awareness: 3), but need some technical expertise to 'weaponize' (Exploitability: 2). With a total risk value of 10, it is classified as a critical risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

4.2.7 V07 - Lateral movement through cluster

The OCP setup will be elaborated on first. Once an attacker can execute commands within a container (Vantage Point: 2) and is able to send network requests (RAL: 2), they can scan the network for other containers, hosts, services, apis or similar interfaces for further access. By default, all containers in all projects (except master & infra components) are put in the same subnet, allowing everyone to communicate with anyone else, including application components belonging to other tenants. These components may rely on the restriction of network access to them instead of authorizing individual requests (Impact: 3). Network scanning tools and accompanying tutorials are readily available (Awareness: 3), but their results depend on the deployed applications. Therefore some technical expertise is needed to leverage the network access (Exploitability: 2). With a total risk value of 7, it is classified as a high risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

4.2.8 V08 - Container breakout

In the OCP setup, a deployed container poses the risk of allowing access to the node it is running on, thus allowing an attacker with the ability to admit new containers (RAL: 2) to "break out" of the container (Vantage Point: 2) and perform actions on the node. This poses a considerable threat, since any container may run on any node by default, allowing an attacker full access to any containers running on the node he controls, which will – especially over time – have a great chance to include containers belonging to other projects (Impact: 3). The OCP default settings limit the possibility of this dramatically, as both privileged containers and those executed as root user are prevented from being run. Network and file system access as well as kernel capabilities are limited, too. In order to affect the host, software vulnerabilities would be needed (Exploitability: 0). The risk lies more in organizations relaxing the default configuration in favor of easy usability, since a majority of container images straight from docker hub require UID 0 and thus a root user. The easiest way to stop those problems this is to permit the default service account within a project access to the 'privileged' SCC

4 Assessing the attack surface risk

permissions. This would significantly increase the risk of a container breakout, but is not considered here as the default configuration is assessed. This is probably the most-talked about attack vector regarding containers (Awareness: 3), but techniques are not obviously documented and breakout methods would have to be customized to the restrictions applied within a cluster. With a total risk value of 5, it is classified as a medium risk.

The AKS setup differs from these circumstances by having no restrictions to admitted pods in place. Thus, this vector is more easily exploitable (Exploitability: 2). The total risk value is increased from 5 to 7 by this, leading to a high risk.

4.2.9 V09 - Image cache compromise

In the OCP setup, a cached container image could be swapped out for a malicious one, thus affecting a container and its application (Impact: 2). In order to do this, an attacker would need administrative access to the underlying node (RAL: 1) and be able to execute commands on it (Vantage Point: 1). This is a sensible vector not commonly discussed (Awareness: 2) and would require technical knowledge to conduct (Exploitability: 2). With a total risk value of 3, it is classified as a low risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

4.2.10 V10 - Container modification at runtime

In the OCP setup, an attacker can try to modify and use an already running container instead of deploying a container with malicious contents,. In order to do this, they would need to execute commands from within the container (Vantage Point: 2) and be able to write to its file system (RAL: 2). This would affect the container accessed and in turn the application it is part of (Impact: 2). Trying to do this within a container is straightforward (Awareness: 3) and would require the same technical expertise as any command line interaction with a linux system (Exploitability: 2). With a total risk value of 5, it is classified as a medium risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

4.2.11 V11 - Resource hoarding (sabotage)

First, the OCP setup will be assessed. With enough access or restrictions too lax, an attacker with access to the cluster interfaces (Vantage Point: 3) and the ability to change or add k8s objects (RAL: 2) may be able to seriously halt the availability of all workloads processed by the cluster by misconfiguration, conducting Denial of Service attacks or wiping nodes or cluster configurations. Since clusters are a complex distributed system with a multitude of configurations, finding the sabotaged component could

take considerable expertise and time if done well, increasing the impact – especially in on-premise environments, where resources are limited. The potential to do this is common sense (Awareness: 3), sabotaging the cluster in a complex and effective way (Impact: 3) may take deeper knowledge and would have to be customized to the environment at hand (Exploitability: 2). With a total risk value of 8, it is classified as a high risk.

The circumstances in the AKS setup differ again by the interfaces being available from the public internet (Vantage Point: 4). Since it is far easier to spin up additional resources in the cloud to continue operations, less impact is estimated (Impact: 2). All other circumstances remain (RAL: 2, Awareness: 3, Exploitability: 2). With a total risk value of 6, it is classified as a medium risk.

4.2.12 V12 - Resource misuse (cryptojacking)

The OCP setup will be assessed first. In order to achieve more monetary gain, an attacker with access to the cluster interfaces (Vantage Point: 3) and the ability to change or add k8s objects (RAL: 2) will try to be discrete. The goal is to (ab)use the computing resources not belonging to and payed for by him to achieve monetary gain though mining cryptocurrencies. Cryptojacking is regularly cited as an up-and-coming attack (Awareness: 3), but to do it without being detected needs some technical skill (Exploitability: 2). If successful for longer periods of time, a companies on-premise resources will be available to a lesser degree as some of them are used for mining. In addition to that, the underlying hardware will be degraded at a higher rate, leading to hardware failures and replacements sooner than normal (Impact: 2). With a total risk value of 5, it is classified as a medium risk.

The circumstances in the AKS setup differ again by the interfaces being available from the public internet (Vantage Point: 4). In contrast to V11, the ability to spin up additional resources in the cloud increases the impact here, since those are virtually unlimited and simply billed towards the attack victim (Impact: 3). All other circumstances remain (RAL: 2, Awareness: 3, Exploitability: 2). With a total risk value of 8, it is classified as a high risk.

4.2.13 V13 - Adding rogue containers

The OCP setup will be assessed first. Instead of manipulating running containers, an attacker with access to the API (Vantage Point: 3) and permissions to spin up containers (RAL: 2) may start their own ones. This is still restricted by container admission restrictions on the user/project, but they can be leveraged to compromise parts of the application (Impact: 2). Doing this is common sense (Awareness: 3), but as in other vectors some technical skill is required to prepare a malicious container (Exploitability: 2) With a total risk value of 5, it is classified as a medium risk.

4 Assessing the attack surface risk

Synonymous to V01, the circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet. (Vantage Point: 4, RAL: 2, Awareness: 3, Exploitability: 2, Impact: 2). The total risk value is increased from 5 to 6 by this, still classifying it as a medium risk.

4.2.14 V14 - Adding rogue nodes

In the OCP setup, an attacker with access to the API (Vantage Point: 3) could try to add a malicious node to the cluster. Similar to V08, they could gain full access to containers of multiple projects by this (Impact: 3). By design, cluster administrator access is needed to add a node within OCP (RAL: 1). This technique is not talked about that much, but still available in public resources and possible in all clusters (Awareness: 2). Docs are publicly available to add nodes to a cluster, basic linux server administration skills are needed to follow them (Exploitability: 2). With a total risk value of 6, it is classified as a medium risk.

The AKS interfaces are again available from the public internet (Vantage Point: 4). Adding nodes would still be possible here, but since they are automatically provisioned and customization of provisioned VMs is not straightforward, this attack is relatively unorthodox (Awareness: 1). The other circumstances remain synonymous to the OCP setup (RAL: 1, Exploitability: 2, Impact: 3). The total risk value remains unchanged at 6, a medium risk.

4.2.15 V15 - Leveraging bad user practice

The OCP setup will be assessed first. User practices outside of the cluster are not affected by cluster controls. Users entering their credentials in faked login windows or publicizing access tokens through public repositories are relatively common occurrences that could potentially allow anyone finding them (RAL: 4) privileged access within that user context (Impact: 2) if they can reach the cluster interfaces (Vantage Point: 3). There are tools available to do this, but doing this is non-intuitive (Awareness: 2) and using them effectively requires some technical skill (Exploitability: 2). With a total risk value of 6, it is classified as a medium risk.

Synonymous to V01, the circumstances in the AKS setup are identical to OCP with the exception of those interfaces being available from the public internet (Vantage Point: 4, RAL: 4, Awareness: 2, Exploitability: 2, Impact: 2). The total risk value remains identical.

4.2.16 V16 - Leveraging bad infrastructure

The OCP setup will be assessed first. The underlying nodes could allow an attacker easy entry, even if the platform itself is secured. Worst case scenarios include nodes available to the public internet

(Vantage Point: 4) with insecure configuration, allowing any attacker unauthorized node access (RAL: 4). Synonymous to V08, this could lead to multiple compromised applications (Impact: 3). Technical know-how would still be needed to leverage this (Exploitability: 2) and attackers are assumed to be unlikely to look for containerized workloads specifically (Awareness: 2). With a total risk value of 9, it is classified as a high risk.

The circumstances in the AKS setup are identical to OCP, resulting in the same risk values.

4.2.17 V17 - Leveraging bad patch management

Known and unpatched vulnerabilities in any components can lead to any number of problems. For both the AKS and OCP setup, worst case values are assumed for all factors except Exploitability, as technical knowledge is required to leverage them. (Vantage Point: 4, RAL: 4, Awareness: 3, Exploitability: 2, Impact: 3). With a total risk value of 10, it is classified as a critical risk for both setups.

5 Managing the attack surface risk

With this chapter, the OWASP Risk Rating Methodology² is continued. As completing the risk management process would exceed the scope of this thesis, the full process will be described and the risk management steps will be demonstrated through two exemplary vectors.

5.1 Defining procedures and approach

Normally one should start with the vector currently having the highest overall risk. We chose two other vectors, since these were more representative of the environments within scope and concise enough to properly present the process. Securing the underlying physical servers or VMs is not an optimal example representing the process of securing a solution based on orchestrated containers.

TODO would start with highest vector, check possible measures and implement when acceptable/ap-plicable. reassess vector risk and either implement more XOR accept residual risk if still highest. otherwise put back into queue -> pick up now highest vector and check possible measures, ... -> repeat until risks are either all below target threshold or all above are accepted.

- Values might decrease through the implementation of security measures, leading to a lower total value. (If multiple techniques could be used to leverage a vector and only one gets its total risk reduced, the new maximum risk value of that vector becomes the vector value(s))

goal is to reduce values above threshold X to below threshold X by applying security measures. This aims to ensure a basic security level, not something against APT groups / zero-day protection / targeted attacks with a lot of resources and competence. (no online banking, user data of average confidentiality etc)

5.2 Managing the risk of V07 - Lateral movement through cluster

TODO this was done in both AKS and OCP, is possible with both default configs!

5.2.1 Demonstrating the successful attack without security measures

TODO: show demo on OCP and AKS, or just one and put the other in appendix?

TODO: show getting in and damage, refer to ppt for commands, console I/O and screenshots

² Fou19a.

5.2.2 Selecting security measures

TODO: use multitenant networkplugin (OCP) and network policies (k8s) TODO: show implementing measures through commands in ppt

The worst AKS-specific problem with this is the mitigation. This risk is not clearly documented in the setup section of the documentation. If one stumbles upon this information in further sections of the docs after setting up his cluster, he might postpone or deny changing the setting to isolate different projects by default. This is because a full cluster rebuild is needed to change this setting!

5.2.3 Demonstration with implemented security measures

TODO show remediation through demo of unsuccessful attack

5.2.4 Risk reassessment

TODO updated table of risk for ocp AND aks

5.3 Managing the risk of V08 - Container breakout

TODO this was done on AKS and OCP, but needed changes in default OCP settings to work.

5.3.1 Demonstrating the successful attack without security measures

TODO: show demo on OCP and AKS, or just one and put the other in appendix?

TODO: show getting in and damage, refer to ppt for commands, console I/O and screenshots

5.3.2 Selecting security measures

TODO dont allow privileged containers (scc in OCP, podsecuritypolicies in AKS) TODO: show implementing measures through commands in ppt

5.3.3 Demonstration with implemented security measures

TODO show remediation through demo of unsuccessful attack

5.3.4 Risk reassessment revisited

TODO SECOND updated table of risk for ocp AND aks <- this one starts with first updated one

5.4 Continuing the risk management process

TODO would start with highest vector, check possible measures and implement when acceptable/applicable. reassess vector risk and either implement more XOR accept residual risk if still highest. otherwise put back into queue -> pick up now highest vector and check possible measures, ... -> repeat until risks are either all below target threshold or all above are accepted.

6 Conclusion

This chapter aims to provide answers to the remaining questions posed in chapter 1.2.

6.0.1 Comparing on-premise and public cloud

TODO reference table 4.1

TODO explain problem of generically comparing, without going into specific solutions -> just comparing OCP vs. AKS. TODO both have pros and cons, regardless of specific solution picked. Specific solutions have their pros and cons, too. cloud risk may be higher initially (both in the solutions we looked at and in general), but there are measures to reduce that to acceptable level for baseline security. some use cases benefit from buy-iaas-provide-paas, since more control with you as provider. high-security requirements on-prem probably better (and may be needed for compliance). -> our POV both work, would have to be decided on a case-by-case basis as they depend on many factors (integration with other offers or services very valuable for cloud-specific, pre-existing on-prem data centers make on-prem easier).

6.0.2 Summary

TODO summary of thesis process - huge scope, could have been better defined beforehand. risk estimate formula still inst perfect, but time would probably still be better allocated to practical stuff or more research, both for thesis and others trying to follow the process. no use in better knowing which one is the biggest problem if no time left to reduce risk in the end. hope that more sources are published and industry standards crystallize so not everyone has to do that much research on their own.

TODO outlook - stuff will continue to move fast, hopefully with security in mind. we are currently far better with security than i.e. pre k8s-1.8. more security features will be introduced or go stable (refer to future k8s pod sandbox option in YAML files). as mentioned in multi-tenant isolation, cluster federation might help and will introduce new solutions, but probably also new problems. rootless builds, currently through buildah etc, becoming more mainstream will probably be the next big security thing here.

TODO commentary - hope this helps someone and spares some time. am excited for future tools and platform developments.

A Appendix

A.1 AKS versions available on May 3rd, 2019

The screenshot shows a web browser window with the URL <https://docs.microsoft.com/en-us/azure/aks/supported-kubernetes-versions>. The page content includes a table of supported Kubernetes versions and their upgrade paths. The table is as follows:

KubernetesVersion	Upgrades
1.12.5	None available
1.12.4	1.12.5
1.11.7	1.12.4, 1.12.5
1.11.6	1.11.7, 1.12.4, 1.12.5
1.10.12	1.11.6, 1.11.7
1.10.9	1.10.12, 1.11.6, 1.11.7
1.9.11	1.10.9, 1.10.12
1.9.10	1.9.11, 1.10.9, 1.10.12

Below the table, there is a section titled 'FAQ' with the question: 'What happens when a customer upgrades a Kubernetes cluster with a minor version that is not supported?'

On the right side of the screenshot, there is an Azure Cloud Shell terminal window. It shows the command `az aks get-versions --location eastus --output table` and its output, which is a table of supported Kubernetes versions and their upgrade paths. The output is as follows:

KubernetesVersion	Upgrades
1.13.5	None available
1.12.7	1.13.5
1.12.6	1.12.7, 1.13.5
1.11.9	1.12.6, 1.12.7
1.11.8	1.11.9, 1.12.6, 1.12.7
1.10.13	1.11.8, 1.11.9
1.10.12	1.10.13, 1.11.8, 1.11.9
1.9.11	1.10.12, 1.10.13
1.9.10	1.9.11, 1.10.12, 1.10.13

The terminal window also shows the command `az aks get-versions --location centralindia --output table` and its output, which is a table of supported Kubernetes versions and their upgrade paths. The output is as follows:

KubernetesVersion	Upgrades
1.13.5	None available
1.12.7	1.13.5
1.12.6	1.12.7, 1.13.5
1.11.9	1.12.6, 1.12.7
1.11.8	1.11.9, 1.12.6, 1.12.7
1.10.13	1.11.8, 1.11.9
1.10.12	1.10.13, 1.11.8, 1.11.9
1.9.11	1.10.12, 1.10.13
1.9.10	1.9.11, 1.10.12, 1.10.13

The terminal window also shows the command `az aks get-versions --location eastasia --output table` and its output, which is a table of supported Kubernetes versions and their upgrade paths. The output is as follows:

KubernetesVersion	Upgrades
1.13.5	None available
1.12.7	1.13.5
1.12.6	1.12.7, 1.13.5
1.11.9	1.12.6, 1.12.7
1.11.8	1.11.9, 1.12.6, 1.12.7
1.10.13	1.11.8, 1.11.9
1.10.12	1.10.13, 1.11.8, 1.11.9
1.9.11	1.10.12, 1.10.13
1.9.10	1.9.11, 1.10.12, 1.10.13

The terminal window also shows the command `az aks get-versions --location westeurope --output table` and its output, which is a table of supported Kubernetes versions and their upgrade paths. The output is as follows:

KubernetesVersion	Upgrades
1.13.5	None available
1.12.7	1.13.5
1.12.6	1.12.7, 1.13.5
1.11.9	1.12.6, 1.12.7
1.11.8	1.11.9, 1.12.6, 1.12.7
1.10.13	1.11.8, 1.11.9
1.10.12	1.10.13, 1.11.8, 1.11.9
1.9.11	1.10.12, 1.10.13
1.9.10	1.9.11, 1.10.12, 1.10.13

The terminal window also shows the command `date` and its output, which is `Fri May 3 07:17:36 UTC 2019`.

Figure A.1 List of k8s versions available in AKS during the practical part of the thesis on May 3rd, 2019. Screenshot taken by Lukas Grams.

A.2 Security advisory email from Microsoft

Azure Kubernetes Service updates for Linux Kernel (SACK) vulnerabilities (CVE-2019-11477, CVE-2019-11478, CVE-2019-11479)

You're receiving this email because you currently use Microsoft Azure Kubernetes Service.

On Monday June 17th, security researchers announced 3 critical security issues impacting the Linux kernel. These are:

- CVE-2019-11477: SACK Panic
- CVE-2019-11478: SACK Slowness
- CVE-2019-11479: Excess Resource Consumption Due to Low MSS Values

These CVEs have been patched by all major Linux vendors. This means your clusters must be updated to mitigate these security issues.

All Azure Kubernetes Service (AKS) customers running unpatched kernels are potentially vulnerable to these security issues. We recommend all customers verify the running kernel and take action if required to apply these updates.

Canonical issued updated, patched kernels and these updated kernels were made available to the AKS customer base as of 2019-06-19.

AKS clusters using the default configuration were patched as of the automatic update on 2019-06-20 00:00 UTC; however, users must reboot their clusters for the patch to take effect.

Customers with clusters created before Friday, June 28, 2019 should confirm their nodes are updated.

Figure A.2 An email from July 15th, 2019, advising users to reboot their AKS clusters in order to apply security patches. Screenshot taken by Lukas Grams.

A.3 AKS cluster default confs?

TODO? <- maybe add screenshots about aks default config situation, leading to rebuild needed to get netpols?




Figure A.3 Default AKS cluster configuration as of July 3rd, 2019. Screenshot taken by Lukas Grams.

Bibliography

Books

- [RH18] Liz Rice and Michael Hausenblas. *Kubernetes Security - How to Build and Operate Applications Securely in Kubernetes*. O'Reilly Media, Inc., 2018. URL: <https://info.aquasec.com/kubernetes-security> (cit. on pp. 3, 18–20, 23, 27).

Online text-based sources

- [Aut19a] The Kubernetes Authors. *Understanding Kubernetes Objects*. 2019. URL: <https://kubernetes.io/docs/concepts/overview/working-with-objects/kubernetes-objects/> (visited on 07/16/2019) (cit. on pp. ix, 9, 10, 13).
- [Inc18a] Red Hat Inc. *Generally Available today: Red Hat OpenShift Container Platform 3.11 is ready to power enterprise Kubernetes deployments*. 2018. URL: <https://www.redhat.com/en/blog/generally-available-today-red-hat-openshift-container-platform-311-ready-power-enterprise-kubernetes-deployments> (visited on 07/16/2019) (cit. on p. 2).
- [Fou19a] OWASP Foundation. *OWASP Risk Rating Methodology*. 2019. URL: https://www.owasp.org/index.php/OWASP_Risk_Rating_Methodology (visited on 07/16/2019) (cit. on pp. 2, 31, 33, 45).
- [Fou19b] OWASP Foundation. *Threat Modeling Cheat Sheet*. 2019. URL: https://github.com/OWASP/CheatSheetSeries/blob/master/cheatsheets/Threat_Modeling_Cheat_Sheet.md (visited on 07/16/2019) (cit. on p. 2).

- [Irw14] Stephen Irwin. *Creating a Threat Profile for Your Organization*. 2014. URL: <https://www.sans.org/reading-room/whitepapers/threats/creating-threat-profile-organization-35492> (visited on 07/16/2019) (cit. on p. 2).
- [SMS11] Murugiah Souppaya, John Morello, and Karen Scarfone. *NIST Special Publication 800-190 Application Container Security Guide*. 2011. URL: <https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-190.pdf> (visited on 07/16/2019) (cit. on pp. 3, 20, 21, 26).
- [Int19a] Center for Internet Security. *CIS Docker Community Edition Benchmark*. 2019. URL: <https://www.cisecurity.org/benchmark/docker/> (visited on 07/16/2019) (cit. on pp. 3, 26).
- [Int19b] Center for Internet Security. *CIS Kubernetes Benchmark*. 2019. URL: <https://www.cisecurity.org/benchmark/kubernetes/> (visited on 07/16/2019) (cit. on pp. 3, 26).
- [RF19] RedGuard and OWASP Foundation. *Container Security Verification Standard*. 2019. URL: <https://github.com/OWASP/Container-Security-Verification-Standard> (visited on 07/16/2019) (cit. on pp. 3, 26).
- [Fou19c] OWASP Foundation. *Docker Security*. 2019. URL: <https://github.com/OWASP/Docker-Security> (visited on 07/16/2019) (cit. on p. 3).
- [Aut19b] The Kubernetes Authors. *Overview of Cloud Native Security*. 2019. URL: <https://kubernetes.io/docs/concepts/security/overview/> (visited on 07/16/2019) (cit. on p. 3).
- [Inc19a] Red Hat Inc. *OpenShift Container Platform 3.11 Container Security Guide*. 2019. URL: https://access.redhat.com/documentation/en-us/openshift_container_platform/3.11/pdf/container_security_guide/OpenShift_Container_Platform-3.11-Container_Security_Guide-en-US.pdf (visited on 07/16/2019) (cit. on p. 3).
- [Cor18] Microsoft Corporation. *Best practices for cluster security and upgrades in Azure Kubernetes Service (AKS)*. 2018. URL: <https://docs.microsoft.com/en-us/>

azure/aks/operator-best-practices-cluster-security (visited on 07/16/2019) (cit. on p. 3).

- [Wat17] Stephen Watts. *SaaS vs PaaS vs IaaS: What's The Difference and How To Choose*. 2017. URL: <https://www.bmc.com/blogs/saas-vs-paas-vs-iaas-whats-the-difference-and-how-to-choose/> (visited on 07/16/2019) (cit. on p. 5).
- [ST11] National Institute of Standards and Technology. *SP 800-145 The NIST Definition of Cloud Computing*. 2011. URL: <https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-145.pdf> (visited on 07/16/2019) (cit. on p. 6).
- [Cor19a] Microsoft Corporation. *What is PaaS? Platform as a Service*. 2019. URL: <https://azure.microsoft.com/en-us/overview/what-is-paas/> (visited on 07/16/2019) (cit. on p. 6).
- [Cor19b] Microsoft Corporation. *What is PaaS? Platform as a Service*. 2019. URL: <https://docs.microsoft.com/en-us/azure/aks/operator-best-practices-cluster-security?view=azuremgmtbilling-1.1.0-preview%5C> (visited on 07/16/2019) (cit. on p. 6).
- [Inc19b] Docker Inc. *What is a Container?* 2019. URL: <https://www.docker.com/resources/what-container> (visited on 07/16/2019) (cit. on p. 6).
- [Raa18] Mike Raab. *Intro to DockerContainers*. 2018. URL: https://static.rainfocus.com/oracle/oraclecode18/session/1513810380873001uIiU/PF/docker-101-ny_1520531065990001vPkT.pdf (visited on 07/16/2019) (cit. on p. 6).
- [Osn18] Rani Osnat. *A brief history of containers: From the 1970s to 2017*. 2018. URL: <https://blog.aquasec.com/a-brief-history-of-containers-from-1970s-chroot-to-docker-2016> (visited on 07/16/2019) (cit. on p. 6).
- [Aut19c] The Kubernetes Authors. *What is Kubernetes*. 2019. URL: <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/> (visited on 07/16/2019) (cit. on p. 7).

- [Cha17] Nick Chase. *OK, I give up. Is Docker now Moby? And what is LinuxKit?* 2017. URL: <https://www.mirantis.com/blog/ok-i-give-up-is-docker-now-moby-and-what-is-linuxkit/> (visited on 07/16/2019) (cit. on pp. 7, 14).
- [Inc19c] Docker Inc. *About Docker Engine*. 2019. URL: <https://docs.docker.com/engine/> (visited on 07/16/2019) (cit. on p. 7).
- [Inc19d] Red Hat Inc. *Getting Started with Containers*. 2019. URL: https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux_atomic_host/7/html-single/getting_started_with_containers/index (visited on 07/16/2019) (cit. on p. 7).
- [Inc18b] Red Hat Inc. *Containers and Images*. 2018. URL: https://docs.openshift.com/enterprise/3.0/architecture/core_concepts/containers_and_images.html (visited on 07/16/2019) (cit. on p. 7).
- [Inc18c] Docker Inc. *Dockerfile Reference*. 2018. URL: <https://docs.docker.com/engine/reference/builder/> (visited on 07/16/2019) (cit. on p. 8).
- [Fou19d] The Linux Foundation. *Open Container Initiative*. 2019. URL: <https://www.opencontainers.org/> (visited on 07/16/2019) (cit. on p. 8).
- [Lew19a] Ian Lewis. *Container Runtimes Part 2: Anatomy of a Low-Level Container Runtime*. 2019. URL: <https://www.ianlewis.org/en/container-runtimes-part-2-anatomy-low-level-contai> (visited on 07/16/2019) (cit. on p. 8).
- [Lew19b] Ian Lewis. *Container Runtimes Part 3: High-Level Runtimes*. 2019. URL: <https://www.ianlewis.org/en/container-runtimes-part-3-high-level-runtimes> (visited on 07/16/2019) (cit. on p. 8).
- [Aut19d] The Kubernetes Authors. *Container Runtimes Part 3: High-Level Runtimes*. 2019. URL: <https://github.com/kubernetes/cri-api/> (visited on 07/16/2019) (cit. on p. 8).
- [Fou19e] The Linux Foundation. *There are over 80 Certified Kubernetes offerings*. 2019. URL: <https://www.cncf.io/certification/software-conformance/> (visited on 07/16/2019) (cit. on pp. 9, 14).

- [Aut19e] The Kubernetes Authors. *Kubernetes Components*. 2019. URL: <https://kubernetes.io/docs/concepts/overview/components/> (visited on 07/16/2019) (cit. on pp. 9, 20).
- [Aut19f] The Kubernetes Authors. *Pod Overview*. 2019. URL: <https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/> (visited on 07/16/2019) (cit. on p. 11).
- [BC18] Matt Butcher and Karen Chu. *Phippy goes to the Zoo*. 2018. URL: <https://www.cncf.io/wp-content/uploads/2018/12/Phippy-Goes-To-The-Zoo.pdf> (visited on 07/16/2019) (cit. on p. 11).
- [Aut19g] The Kubernetes Authors. *ReplicaSet*. 2019. URL: <https://kubernetes.io/docs/concepts/workloads/controllers/replicaset/> (visited on 07/16/2019) (cit. on p. 11).
- [Aut19h] The Kubernetes Authors. *Picking the Right Solution*. 2019. URL: <https://v1-13.docs.kubernetes.io/docs/setup/pick-right-solution/#> (visited on 07/16/2019) (cit. on p. 13).
- [Aut19i] The Kubernetes Authors. *The Kubernetes API*. 2019. URL: <https://kubernetes.io/docs/concepts/overview/kubernetes-api/> (visited on 07/16/2019) (cit. on p. 13).
- [Inc19e] Red Hat Inc. *OpenShift plans and pricing*. 2019. URL: <https://www.openshift.com/products/pricing/> (visited on 07/16/2019) (cit. on p. 14).
- [Har19] Brian Harrington. *OpenShift and Kubernetes: What's the difference?* 2019. URL: <https://www.redhat.com/en/blog/openshift-and-kubernetes-whats-difference> (visited on 07/16/2019) (cit. on p. 14).
- [Inc19f] Red Hat Inc. *The Origin Community Distribution of Kubernetes that powers Red Hat OpenShift*. 2019. URL: <https://www.okd.io/> (visited on 07/16/2019) (cit. on p. 14).
- [Inc19g] Red Hat Inc. *OpenShift Container Platform 3.x Tested Integrations*. 2019. URL: <https://access.redhat.com/articles/2176281> (visited on 07/16/2019) (cit. on pp. 14, 30).

- [Inc19h] Red Hat Inc. *Get Started with the CLI*. 2019. URL: https://docs.openshift.com/container-platform/3.11/cli_reference/get_started_cli.html (visited on 07/16/2019) (cit. on p. 14).
- [Inc19i] Red Hat Inc. *Managing Projects*. 2019. URL: https://docs.openshift.com/container-platform/3.11/admin_guide/managing_projects.html (visited on 07/16/2019) (cit. on p. 14).
- [Inc19j] Red Hat Inc. *Using the CRI-O Container Engine*. 2019. URL: https://docs.openshift.com/container-platform/3.11/crio/crio_runtime.html (visited on 07/16/2019) (cit. on p. 14).
- [Cor19c] Microsoft Corporation. *Integrate Azure Active Directory with Azure Kubernetes Service*. 2019. URL: <https://docs.microsoft.com/en-us/azure/aks/azure-ad-integration> (visited on 07/16/2019) (cit. on p. 15).
- [Aut19j] The Kubernetes Authors. *Kubernetes version and version skew support policy*. 2019. URL: <https://kubernetes.io/docs/setup/release/version-skew-policy/> (visited on 07/16/2019) (cit. on pp. 17, 27).
- [Aut19k] The Kubernetes Authors. *Web UI (Dashboard)*. 2019. URL: <https://kubernetes.io/docs/tasks/access-application-cluster/web-ui-dashboard/> (visited on 07/17/2019) (cit. on p. 18).
- [Inc19k] Red Hat Inc. *Create and Build an Image Using the Web Console*. 2019. URL: https://docs.openshift.com/container-platform/3.11/getting_started/developers_console.html (visited on 07/17/2019) (cit. on p. 18).
- [Cor19d] Microsoft Corporation. *Quickstart: Deploy an Azure Kubernetes Service (AKS) cluster using the Azure portal*. 2019. URL: <https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-portal> (visited on 07/17/2019) (cit. on p. 18).
- [Gom18] Paulo Gomez. *How a naughty docker image on AKS could give an attacker access to your Azure Subscription?* 2018. URL: <https://itnext.io/how-a-naughty->

docker-image-on-aks-could-give-an-attacker-access-to-your-azure-subscription-6d05b92bf811 (visited on 07/17/2019) (cit. on p. 18).

- [Aut19l] The Kubernetes Authors. *Network Policies*. 2019. URL: <https://kubernetes.io/docs/concepts/services-networking/network-policies/> (visited on 07/17/2019) (cit. on p. 19).
- [Aut19m] The Kubernetes Authors. *Auditing*. 2019. URL: <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/> (visited on 07/17/2019) (cit. on p. 19).
- [San17] Jorge Salamero Sanz. *Monitoring Kubernetes (part 2): Best practices for alerting on Kubernetes*. 2017. URL: <https://sysdig.com/blog/alerting-kubernetes/> (visited on 07/17/2019) (cit. on p. 19).
- [Aut19n] The Kubernetes Authors. *Pod Security Policies*. 2019. URL: <https://kubernetes.io/docs/concepts/policy/pod-security-policy/> (visited on 07/17/2019) (cit. on p. 20).
- [MI18] Jeff Morgan and Docker Inc. *Introducing the New Docker Hub*. 2018. URL: <https://blog.docker.com/2018/12/the-new-docker-hub/> (visited on 07/17/2019) (cit. on p. 21).
- [RH19] Liz Rice and Michael Hausenblas. *Kubernetes Security*. 2019. URL: <https://kubernetes-security.info/> (visited on 07/17/2019) (cit. on p. 21).
- [con19] Open Policy Agent contributors. *Kubernetes Admission Control*. 2019. URL: <https://kubernetes-security.info/> (visited on 07/18/2019) (cit. on p. 22).
- [Aut19o] The Kubernetes Authors. *Network Policies*. 2019. URL: <https://kubernetes.io/docs/concepts/services-networking/network-policies/> (visited on 07/18/2019) (cit. on p. 22).
- [Aut19p] Istio Authors. *What is Istio?* 2019. URL: <https://istio.io/docs/concepts/what-is-istio/> (visited on 07/18/2019) (cit. on p. 22).

- [Wal14] Daniel J Walsh. *Are Docker containers really secure?* 2014. URL: <https://opensource.com/business/14/7/docker-security-selinux> (visited on 07/18/2019) (cit. on p. 22).
- [Fou19f] OWASP Foundation. *Addition to the threat mindmap might be needed.* 2019. URL: <https://github.com/OWASP/Docker-Security/issues/9> (visited on 07/18/2019) (cit. on p. 22).
- [CB19] Dominik Czarnota and Trail of Bits. *Understanding Docker container escapes.* 2019. URL: <https://blog.trailofbits.com/2019/07/19/understanding-docker-container-escapes/> (visited on 07/29/2019) (cit. on p. 22).
- [Inc19l] Docker Inc. *Content trust in Docker.* 2019. URL: https://docs.docker.com/engine/security/trust/content_trust/ (visited on 07/18/2019) (cit. on p. 23).
- [TK18] Jen Tong and Maya Kaczorowski. *Kubernetes Runtime Security.* 2018. URL: https://static.sched.com/hosted_files/kccnceu18/72/20180524%20-%20KubeCon%20EU%20-%20Kubernetes%20Runtime%20Security.pdf (visited on 07/18/2019) (cit. on p. 23).
- [Aut19q] The Kubernetes Authors. *Resource Quotas.* 2019. URL: <https://kubernetes.io/docs/concepts/policy/resource-quotas/> (visited on 07/19/2019) (cit. on p. 24).
- [Aut19r] The Kubernetes Authors. *Tools for Monitoring Resources.* 2019. URL: <https://kubernetes.io/docs/tasks/debug-application-cluster/resource-usage-monitoring/> (visited on 07/19/2019) (cit. on p. 24).
- [Tea18] RedLock CSI Team. *Lessons from the Cryptojacking Attack at Tesla.* 2018. URL: <https://redlock.io/blog/cryptojacking-tesla> (visited on 07/19/2019) (cit. on p. 24).
- [Aut19s] The Kubernetes Authors. *Configure Access to Multiple Clusters.* 2019. URL: <https://kubernetes.io/docs/tasks/access-application-cluster/>

configure-access-multiple-clusters/ (visited on 07/19/2019) (cit. on p. 25).

- [EdO19] EdOverflow. *“CI Knew There Would Be Bugs Here” — Exploring Continuous Integration Services as a Bug Bounty Hunter*. 2019. URL: <https://edoverflow.com/2019/ci-knew-there-would-be-bugs-here/> (visited on 07/19/2019) (cit. on p. 25).
- [SN19] Vitaly Simonovich and Ori Nakar. *Hundreds of Vulnerable Docker Hosts Exploited by Cryptocurrency Miners*. 2019. URL: <https://www.imperva.com/blog/hundreds-of-vulnerable-docker-hosts-exploited-by-cryptocurrency-miners/> (visited on 07/21/2019) (cit. on p. 26).
- [Tec19] Graz University of Technology. *Meltdown and Spectre*. 2019. URL: <https://meltdownattack.com/> (visited on 07/21/2019) (cit. on p. 26).
- [Int19c] Center for Internet Security. *CIS Benchmarks*. 2019. URL: <https://www.cisecurity.org/cis-benchmarks/> (visited on 07/21/2019) (cit. on p. 26).
- [Sco19] Stuart Scott. *AWS Security: Bastion Host, NAT instances and VPC Peering*. 2019. URL: <https://cloudacademy.com/blog/aws-bastion-host-nat-instances-vpc-peering-security/> (visited on 07/21/2019) (cit. on p. 26).
- [Ked19] Mark Kedgley. *The Problem with Running Outdated Software*. 2019. URL: <https://www.newnettechnologies.com/whitepaper/Outdated-Software-Whitepaper.pdf> (visited on 07/21/2019) (cit. on p. 26).
- [Dio19] Yuri Diogenes. *Using Azure Monitor to send an Email Notification for Azure Security Center Alerts*. 2019. URL: <https://blogs.technet.microsoft.com/yuridiogenes/2018/08/01/using-azure-monitor-to-send-an-email-notification-for-azure-security-center-alerts/> (visited on 07/21/2019) (cit. on p. 27).
- [Inc19m] Red Hat Inc. *Red Hat OpenShift Container Platform Life Cycle Policy*. 2019. URL: <https://access.redhat.com/support/policy/updates/openshift> (visited on 07/21/2019) (cit. on p. 27).

- [Inc19n] Red Hat Inc. *Planning your installation*. 2019. URL: <https://docs.openshift.com/container-platform/3.11/install/index.html> (visited on 07/30/2019) (cit. on p. 30).
- [FIR19] FIRST.Org. *Common Vulnerability Scoring System version 3.1: User Guide*. 2019. URL: <https://www.first.org/cvss/user-guide> (visited on 07/29/2019) (cit. on p. 31).

Online video sources

- [McC18] Rory McCune. *A Hacker's Guide to Kubernetes and the Cloud - Rory McCune, NCC Group PLC (Intermediate Skill Level)*. 2018. URL: <https://youtu.be/dxKpCO2dAy8?t=1758> (visited on 04/25/2019) (cit. on p. 4).
- [Ric19] Liz Rice. *DIY Pen-Testing for Your Kubernetes Cluster - Liz Rice, Aqua Security*. 2019. URL: <https://www.youtube.com/watch?v=fVqCAUJiIn0&feature=youtu.be&t=525> (visited on 07/17/2019) (cit. on p. 18).
- [Gee17] Brad Geesaman. *Hacking and Hardening Kubernetes Clusters by Example [I] - Brad Geesaman, Symantec*. 2017. URL: <https://www.youtube.com/watch?v=vTgQLzeBfRU&feature=youtu.be&t=185> (visited on 07/16/2019) (cit. on p. 29).

Other sources

- [Mei19] Nico Meisenzahl. *docker-drawings*. A collection of graphics given to Lukas Grams by Nico Meisenzahl at the Global Azure Bootcamp Rosenheim on 2019-04-27. 2019 (cit. on pp. 10–12).